

# Retrieval-Augmented Generation

---

Lara J. Martin (she/they)

<https://laramartin.net/interactive-fiction-class>

*Slides from ACL 2023 Tutorial by Akari Asai, Sewon Min, Zexuan Zhong, & Dr. Danqi Chen*

# Learning Objectives

---

Define the Story Cloze Test and determine its place in guided story generation

Understand the reasons why RAG was created

Explore how the retrieval component interacts with the LLM in RAG

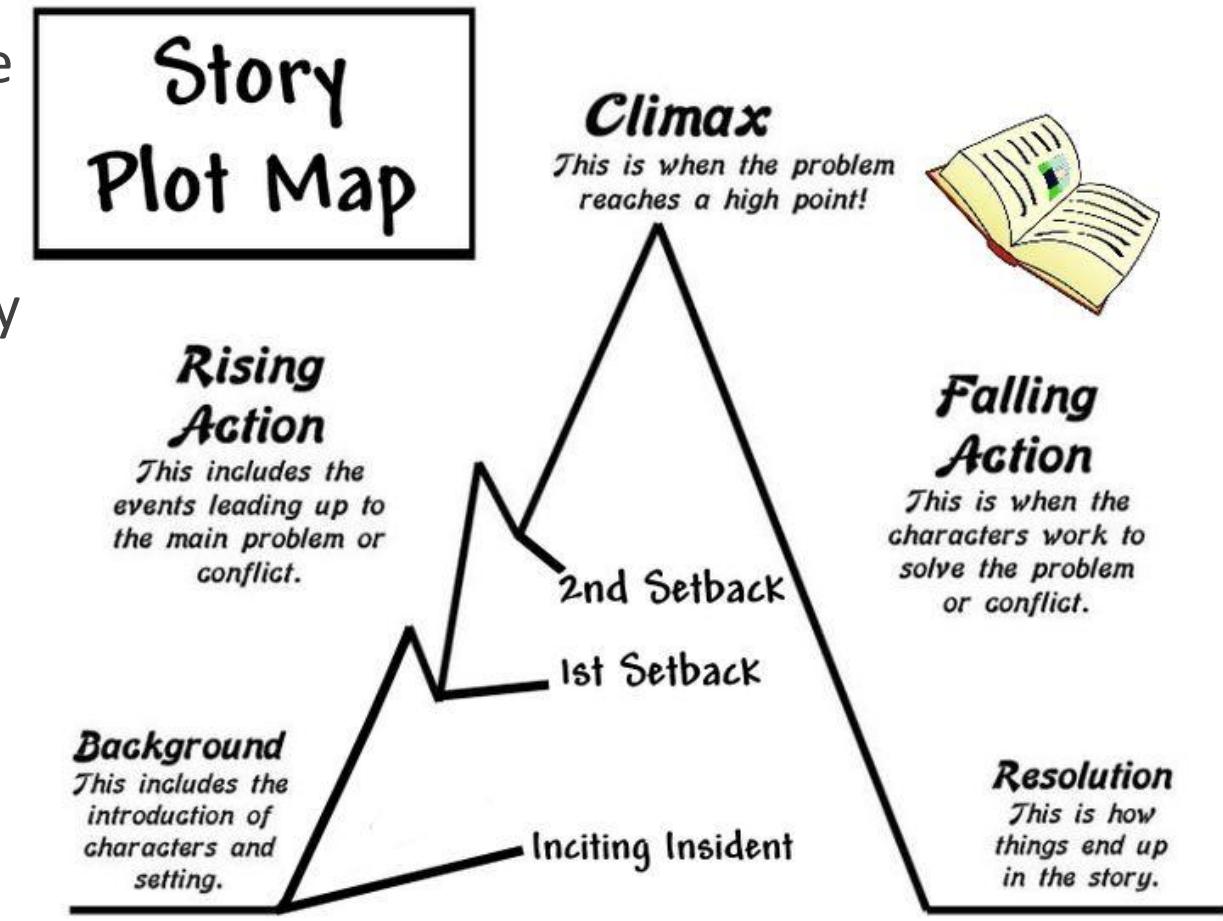
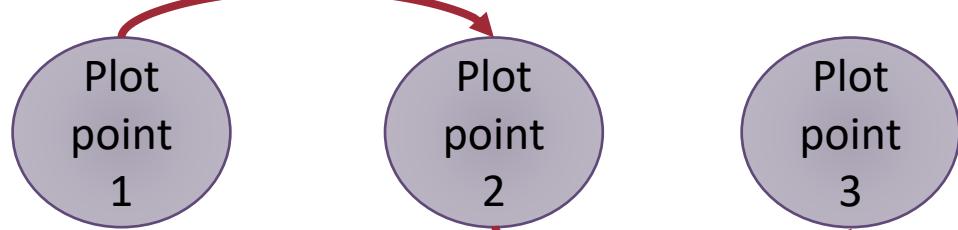
Extract implementation details from papers and find different ways RAG is implemented

Compare plot-guided generation to retrieval-augmented generation for stories

# Review: Scripts, Procedures, and Plots...oh my!

Schank & Abelson believe that everyone has scripts in their heads built from common experiences

Authors often plan out **plots** before they write stories



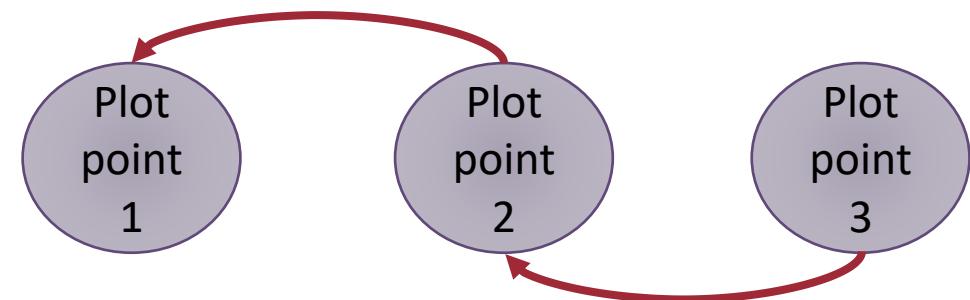
# Review: Scripts, Procedures, and Plots...oh my!

---

Schank & Abelson believe that everyone has scripts in their heads built from common experiences

Authors often plan out plots before they write stories

Stories that aren't planned out either have to "**reincorporate**"<sup>[1]</sup> ideas or the stories feel unfinished



[1] The idea of *reincorporation* is explored in the book [Impro by Keith Johnstone](#)

# Review: Ways of Extracting Plot Points

Most salient keywords

Event representations

Verb-Noun Sets

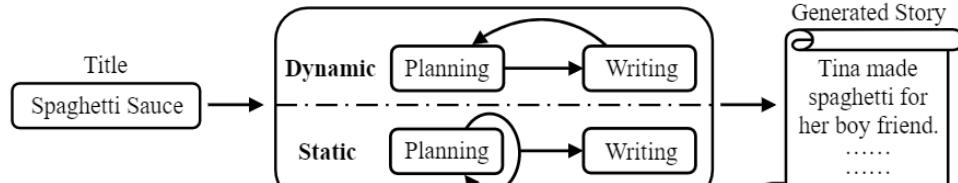
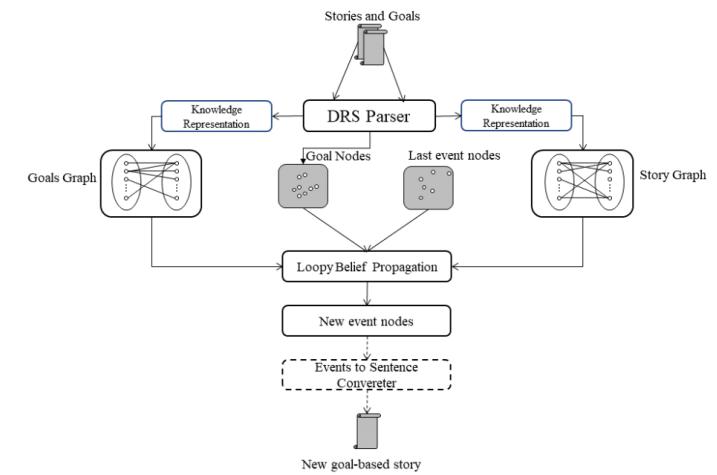
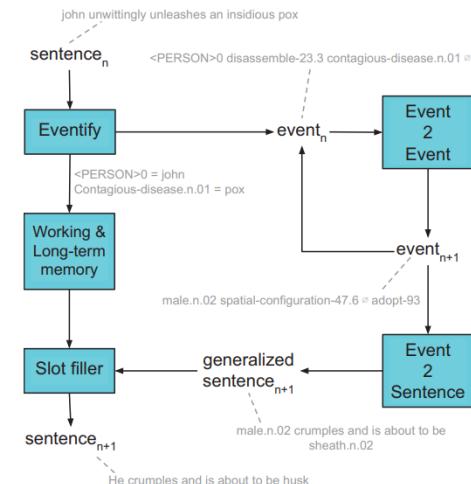


Figure 1: An overview of our system.

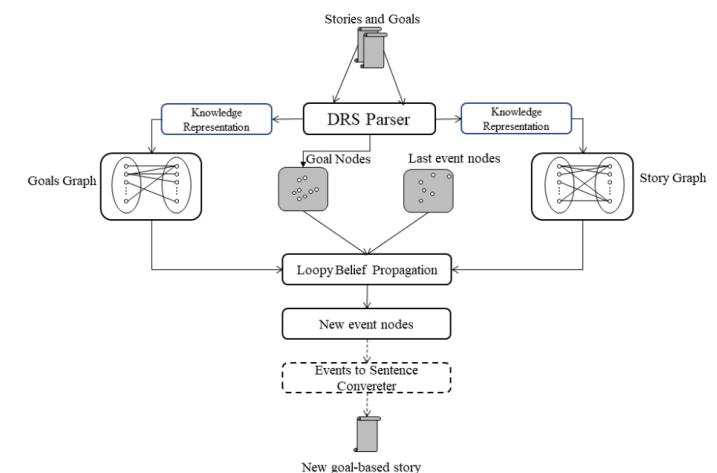
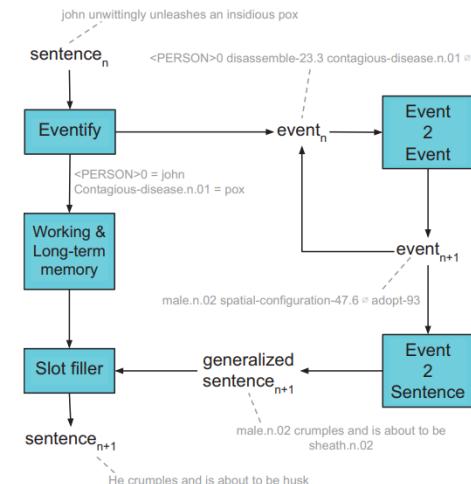
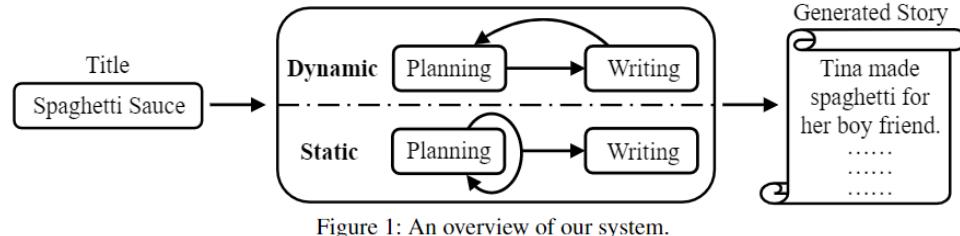


# Review: Generating with Plot Points

Co-generated vs conditioned (prompted) with plot

Generate event & then translate to natural language

Graph algorithms (Loopy Belief Propagation)



# The Story Cloze Test

---

# What is a Cloze Test?

---

- Something is removed from a text; try to guess what's missing
- Used for reading comprehension, grammar, etc. (with humans)

# Unsupervised Learning of Narrative Event Chains

Nathanael Chambers and Dan Jurafsky

Department of Computer Science

Stanford University

Stanford, CA 94305

{natec, jurafsky}@stanford.edu

## Abstract

Hand-coded *scripts* were used in the 1970-80s as knowledge backbones that enabled inference and other NLP tasks requiring deep semantic knowledge. We propose unsupervised induction of similar schemata called *narrative event chains* from raw newswire text.

A narrative event chain is a partially ordered set of events related by a common protagonist. We describe a three step process to learning narrative event chains. The first uses unsupervised distributional methods to learn narrative relations between events sharing coreferencing arguments. The second applies a tem-

tate learning, and thus this paper addresses the three tasks of chain induction: *narrative event induction*, *temporal ordering of events* and *structured selection* (pruning the event space into discrete sets).

Learning these prototypical schematic sequences of events is important for rich understanding of text. Scripts were central to natural language understanding research in the 1970s and 1980s for proposed tasks such as summarization, coreference resolution and question answering. For example, Schank and Abelson (1977) proposed that understanding text about restaurants required knowledge about the Restaurant Script, including the participants (Customer, Waiter, Cook, Tables, etc.), the events (order, pay, etc.)

# Narrative Cloze Test

---

Evaluate “event relatedness”

Find which events could be missing from a narrative chain

Uses verbs only

# Narrative Cloze Test

---

**Known events:**

(pleaded subj), (admits subj), (convicted obj)

**Likely Events:**

sentenced obj	0.89	indicted obj	0.74
paroled obj	0.76	fined obj	0.73
fired obj	0.75	denied subj	0.73

X pleaded \_  
X admits \_  
\_ convicted X

Figure 1: Three narrative events and the six most likely events to include in the same chain.

# Finish the story

---

Gina was worried the cookie dough in the tube would be gross.

She was very happy to find she was wrong.

The cookies from the tube were as good as from scratch.

Gina intended to only eat 2 cookies and save the rest.

- A. Gina liked the cookies so much she ate them all in one sitting. 
- B. Gina gave the cookies away at her church.

# Story Cloze Test

---

Predict/select the most likely story \*ending\*

- Given the first 4 sentences of the story

Full sentences

Multiple choice evaluation

## An RNN-based Binary Classifier for the Story Cloze Test

Melissa Roemmle<sup>\*</sup>  
Institute for Creative Technologies  
University of Southern California  
roemmle@ict.usc.edu

Naoya Inoue  
Tohoku University  
naoya-i@ecei.tohoku.ac.jp

Sosuke Kobayashi<sup>\*</sup>  
Preferred Networks, Inc.  
sosk@preferred.jp

Andrew M. Gordon  
Institute for Creative Technologies  
University of Southern California  
gordon@ict.usc.edu

Poster Presentation

CIKM '20, October 19–23, 2020, Virtual Event, Ireland

## Toward Better Storylines with Sentence-Level Language Models

Daphne Ippolito<sup>\*</sup>  
daphnei@seas.upenn.edu

David Grangier  
grangier@google.com

Douglas Eck  
deck@google.com

Chris Callison-Burch  
ccb@seas.upenn.edu

### Abstract

We propose a sentence-level language model which selects the next sentence in a story from a finite set of fluent alternatives. Since it does not need to model fluency, the sentence-level language model can focus on longer range dependencies, which are crucial for multi-sentence coherence. Rather than dealing with individual words, our method treats the story so far as a list of pre-trained sentence embeddings and predicts an embedding for the next sentence, which is more efficient than predicting word embeddings. Notably this allows us to consider a large number of candidates for the next sentence during training. We demonstrate the effectiveness of our approach with state-of-the-art accuracy on the unsupervised Story Cloze task and with promising results on larger-scale next sentence prediction tasks.

sequence of imagined roles (Liu et al., 2019).

Our work is the first to consider how to pose a model which can handle long-range dependencies of context and a large set of fluent alternatives. We use pre-trained sentence embeddings (Devlin et al., 2019) to build a model which can handle long-range dependencies of context and a large set of fluent alternatives. Given the emphasis of the story, our model uses a sentence embedding of the story, which is more efficient than predicting word embeddings. Notably this allows us to consider a large number of candidates for the next sentence during training. We demonstrate the effectiveness of our approach with state-of-the-art accuracy on the unsupervised Story Cloze task and with promising results on larger-scale next sentence prediction tasks.

This task is to predict the next sentence in a story given the previous sentences. Our model only considers the last sentence in the story. This is because the previous sentences provide enough context for the model to predict the next sentence. The model then generates a sentence embedding for the next sentence based on the previous sentence's embedding and the current sentence's embedding. The model then compares the generated sentence embedding with the sentence embeddings of all the other sentences in the story. The sentence with the highest similarity is selected as the next sentence.

## Tackling the Story Ending Biases in The Story Cloze Test

Rishi Sharma<sup>1</sup>, James F. Allen<sup>1,2</sup>, Omid Bakhshandeh<sup>3</sup>, Nasrin Mostafazadeh<sup>4\*</sup>

<sup>1</sup> University of Rochester, <sup>2</sup> Institute for Human and Machine Cognition, <sup>3</sup> Verneek.ai <sup>4</sup> Elemental Cognition  
rishi.sharma@rochester.edu, nasrinm@cs.rochester.edu

### Abstract

The Story Cloze Test (SCT) is a recent framework for evaluating story comprehension and script learning. There have been a variety of models tackling the SCT so far. Although the original goal behind the SCT was to require systems to perform deep language understanding and commonsense reasoning for successful narrative understanding, some recent models could perform significantly better than the initial baselines by leveraging human-authorship biases discovered in the SCT dataset. In order to shed some

light on this issue, we propose a new test, called the Multi-Story Cloze Test (MSCT), which provides a more challenging evaluation of story comprehension. MSCT requires systems to handle multiple stories simultaneously, and to reason about the relationships between them. This test evaluates a story comprehension system where the system is given a four-sentence short story as the ‘context’ and two alternative endings and to the story, labeled ‘right ending’ and ‘wrong ending.’ Then, the system’s task is to choose the right ending. In order to support this task, Mostafazadeh et al. also provide the ROC Stories dataset, which is a collection of crowd-sourced complete five sentence stories through Amazon Mechanical Turk (MTurk). Each story follows a character through a fairly simple series of events to a conclusion.

Several shallow and neural models, including the state-of-the-art script learning approaches, were presented as baselines (Mostafazadeh et al., 2019; Allen et al., 2019; Liu et al., 2019).

## Enhanced Story Representation by ConceptNet for Predicting Story Endings

Shanshan Huang  
huangss\_33@sjtu.edu.cn  
Shanghai Jiao Tong University

Libin Shen  
libin@leyantech.com  
Leyan Tech

Kenny Q. Zhu<sup>\*</sup>  
kzhu@cs.sjtu.edu.cn  
Shanghai Jiao Tong University

Qianzi Liao  
liaoqz@sjtu.edu.cn  
Shanghai Jiao Tong University

Yinggong Zhao  
ygzhaoo@leyantech.com  
Leyan Tech

### ABSTRACT

Predicting endings for machine commonsense representation of the story is a challenging task. Pre-trained language models have shown success in this task by exploiting the Story Cloze dataset instead of “unseen” endings. In this paper, we propose to improve the sentences to predict the endings by exploring the latent relationship between the endings and the enhanced sentence representations. Our proposed language models, makes the most of the popular Story Cloze data.

### CCS CONCEPTS

IEEE/ACM TRANSACTIONS ON AUDIO, SPEECH, AND LANGUAGE PROCESSING, VOL. 27, NO. 4, APRIL 2019

719

## Story Ending Selection by Finding Hints From Pairwise Candidate Endings

Mantong Zhou<sup>\*</sup>, Minlie Huang<sup>\*</sup>, and Xiaoyan Zhu

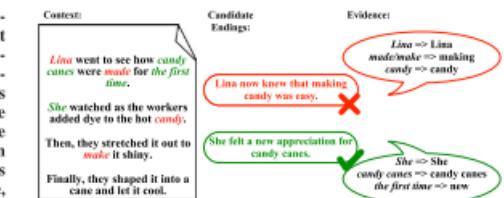


Fig. 1. Evidence bias issue: both a wrong ending (in red) and a correct ending (in green) can obtain sufficient evidence from the story context.

important linkages between a story context and a candidate ending. They suffer from the issue of **evidence bias**: both the wrong and correct endings can obtain sufficient support from the story context. As illustrated in Fig. 1, the wrong ending (in red) and the correct ending (in green) can be supported by the red-colored evidence and the green-colored evidence in the story context, respectively. Thus, it is difficult for matching-based models to distinguish such cases. The situation is not rare because both correct and wrong endings are written to fit the world of a story.

# Think Pair Share

---

The Story Cloze Test was created for evaluating systems' performance on understanding stories.

How could you use it instead for *generation*?

# Retrieval-Augmented Generation

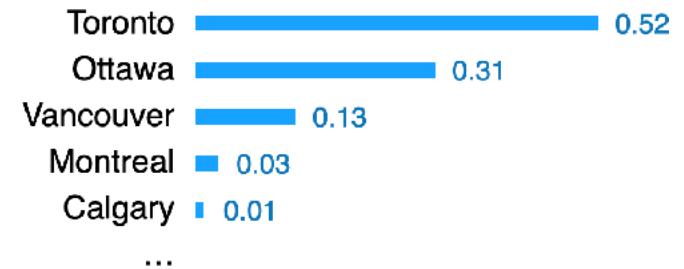
---

Retrieval-based language models (LMs)  
Retrieval-based LMs = Retrieval + LMs

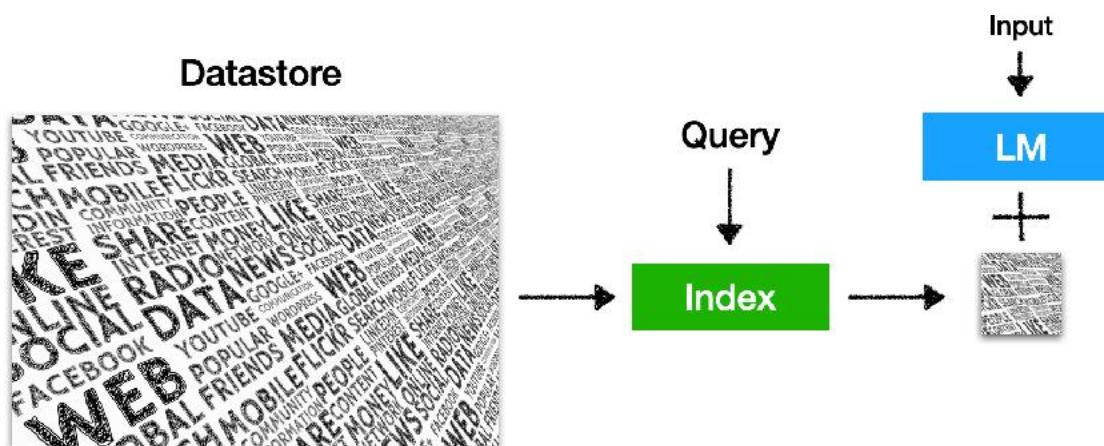
- It is a **language model**  $P(x_n \mid x_1, x_2, \dots, x_{n-1})$

The capital city of Ontario is \_\_\_\_\_

(can be broadly extended to masked language models or encoder-decoder models)



- It retrieves from an **external datastore** (at least during inference time)



# Retrieval for knowledge-intensive NLP tasks

**Representative tasks:** open-domain QA, fact checking, entity linking, ..



Image: <http://ai.stanford.edu/blog/retrieval-based-NLP/>

Drives a lot of research on better algorithms for **dense retrieval**, e.g., **DPR** (Karpukhin et al., 2020), **CoBERT** (Khattab and Zaharia, 2020), **ANCE** (Xiong et al., 2021), **Contriever** (Izacard et al., 2022), ...

Why retrieval LMs?

# Why retrieval-based LMs?



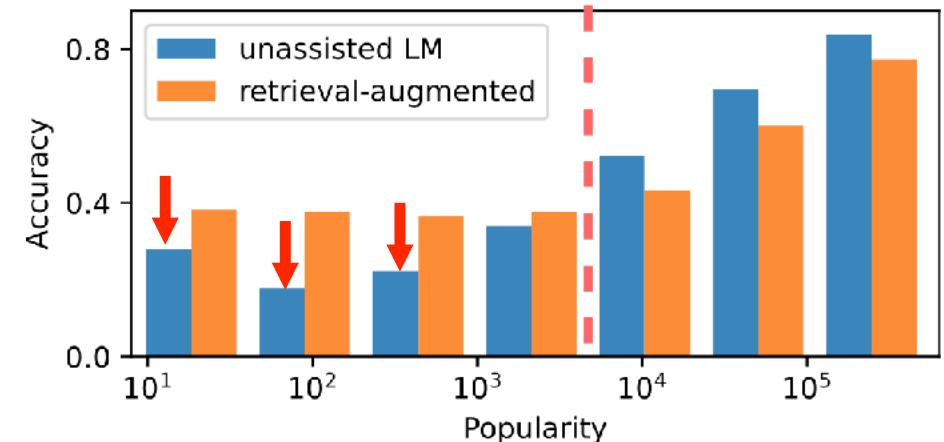
List 5 important papers authored by Geoffrey Hinton



Geoffrey Hinton is a renowned computer scientist ...  
Here are five important papers authored by him:

1. **"Learning Internal Representations by Error Propagation"** (with D. E. Rumelhart and R. J. Williams) - This paper, published in **1986**, ...
2. **"Deep Boltzmann Machines"** (with R. Salakhutdinov) - Published in **2009**, ...
3. **"Deep Learning"** (with Y. Bengio and A. Courville) - Published as a book in **2016**, ...
4. **"Attention Is All You Need"** (with V. Vaswani, N. Shazeer, et al.) - Published in **2017**, this paper introduced the Transformer model, ...

What is Kathy Saltzman's occupation?



(Mallen et al., 2023)

GPT-3 davinci-003: 20%-30% accuracy

## Why retrieval-based LMs?

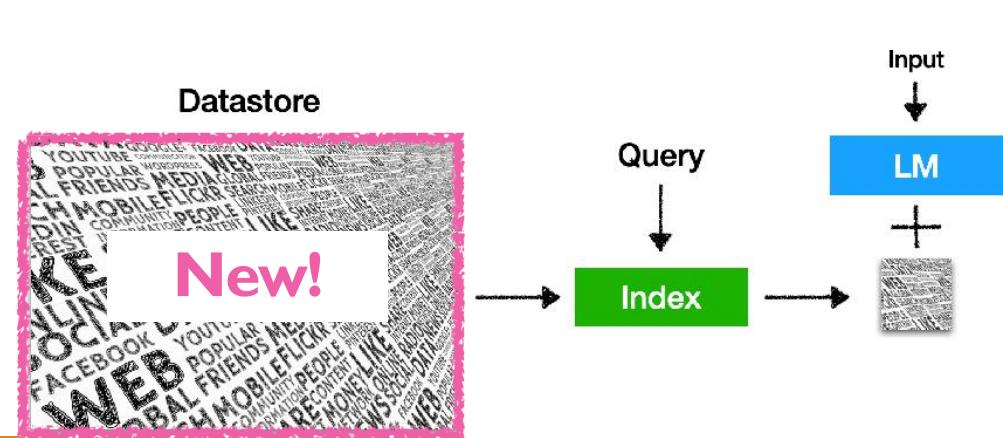
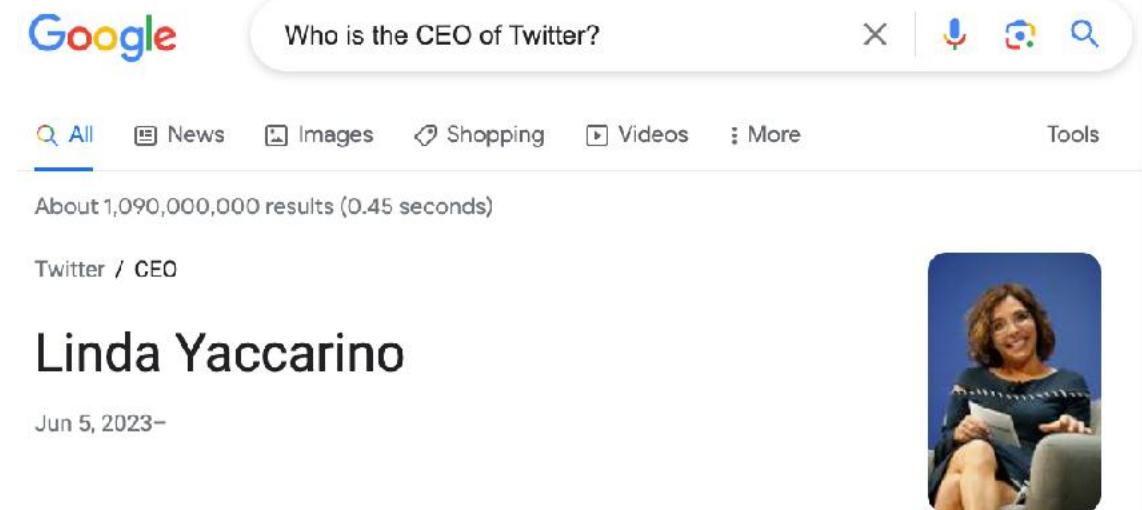


Who is the CEO of Twitter?



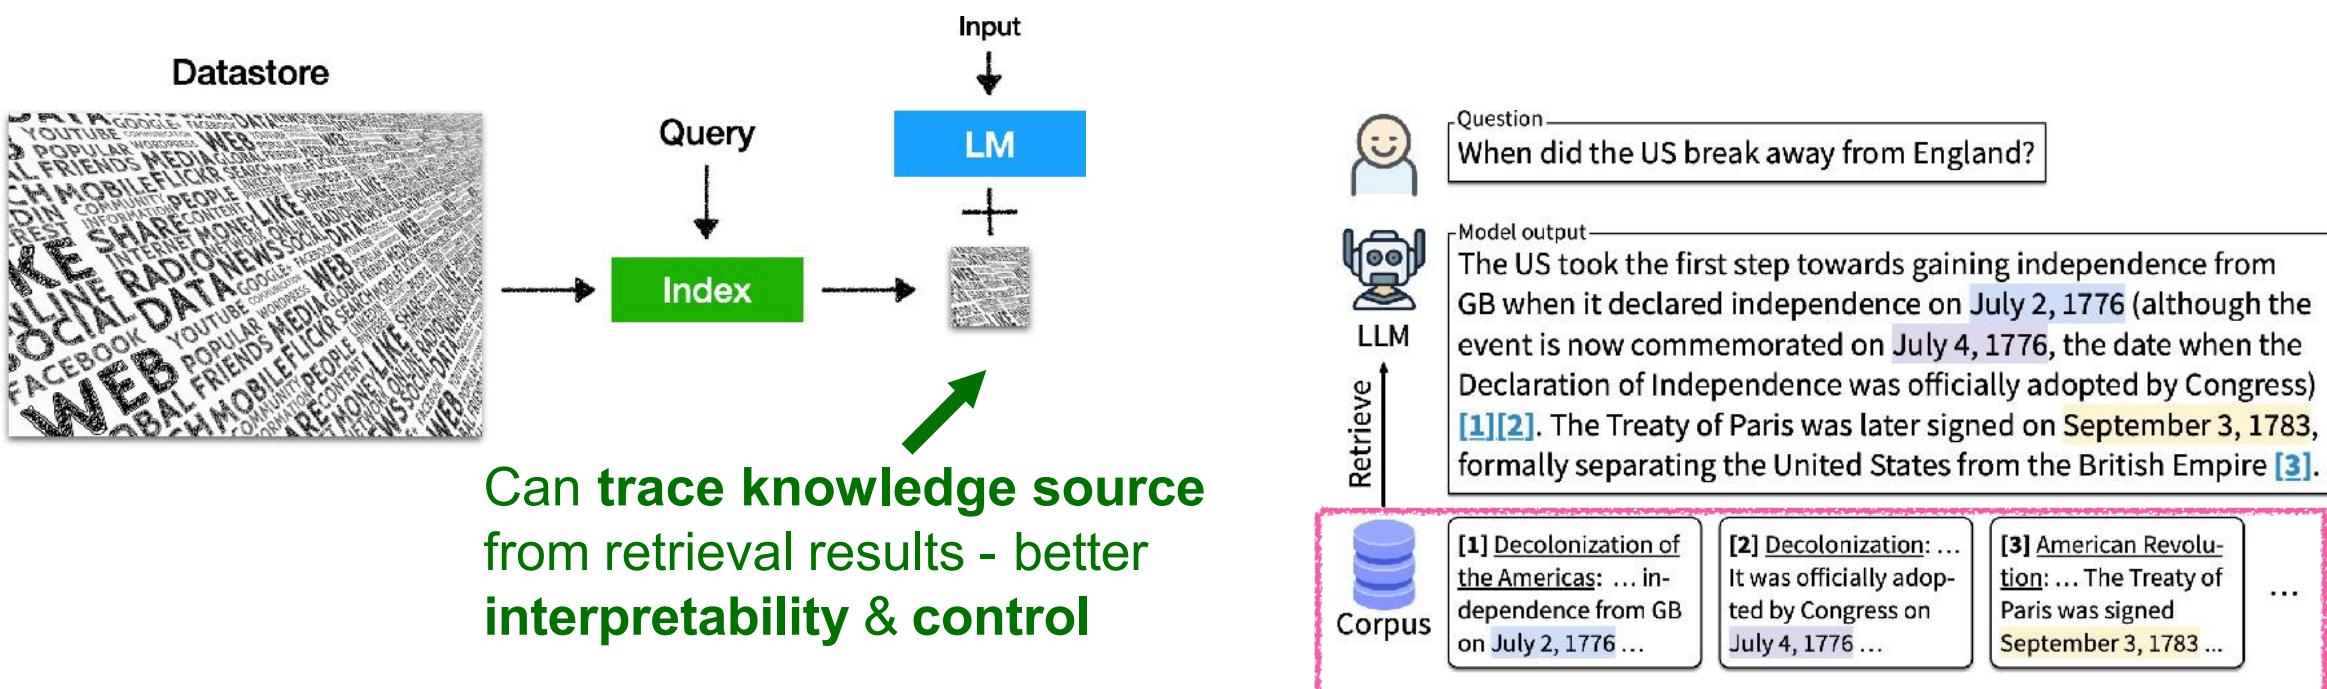
As of my **knowledge cutoff in September 2021**, the CEO of Twitter is **Jack Dorsey**....

- Existing **knowledge editing** methods are still NOT scalable (**active research!**)
- The datastore can be easily **updated** and **expanded** - even without retraining!



# Why retrieval-based LMs?

Generating text with citations



(Nakano et al. 2021; Menick et al., 2022; Gao et al., 2023)

# Why retrieval-based LMs?

Create an itinerary for exploring Toronto over a span of two days.

Toronto is a vibrant and diverse city with many attractions and activities to suit different tastes and preferences. Here are some suggestions for creating an itinerary for exploring Toronto over two days:

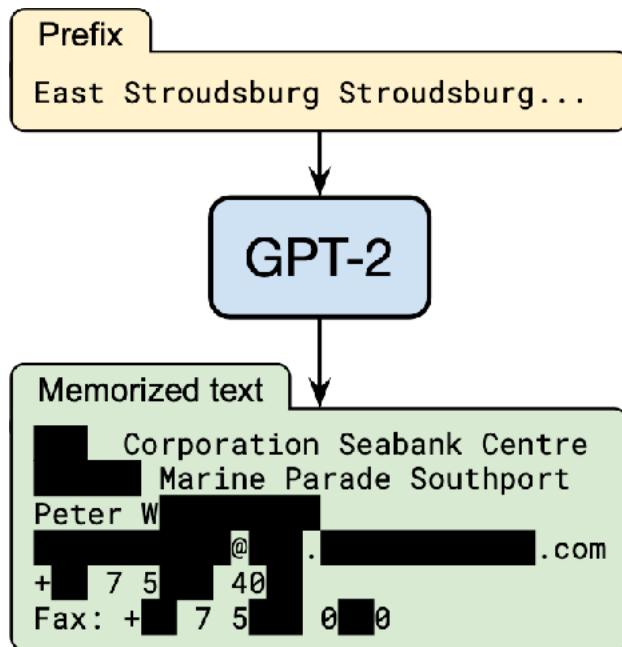
- Day 1: Start your day with a visit to the **CN Tower**, the iconic landmark that offers panoramic views of the city and beyond <sup>1</sup>. You can also experience the thrill of walking on the glass floor or the EdgeWalk, the world's highest hands-free walk <sup>1</sup>. Next, head to the nearby **St. Lawrence Market**, one of the world's best food markets, where you can sample a variety of cuisines and local specialties <sup>2</sup>. After lunch, take a stroll along **Queen West**, a trendy neighborhood with eclectic shops, galleries, cafes and street art <sup>3</sup>. In the evening, enjoy a show at one of the many theaters or comedy clubs in the **Entertainment District**, or catch a game at the **Scotiabank Arena** if you're a sports fan.
- Day 2: Explore the history and culture of Toronto at the **Royal Ontario Museum**, one of the largest and most comprehensive museums in North America, featuring exhibits on art, natural history, world cultures and more <sup>4</sup>. Then, hop on a ferry to the **Toronto Islands**, a group of islands that offer a relaxing escape from the city, with beaches, parks, trails and amusement rides <sup>3</sup> <sup>5</sup>. You can also rent a bike or kayak to explore the islands at your own pace. For dinner, head to **Chinatown**, one of the largest and most vibrant in North America, where you can find a variety of Asian cuisines and shops <sup>3</sup>.

I hope this helps you plan your trip to Toronto. Have fun! 😊

Learn more:

1. [cntower.ca](http://cntower.ca)
2. [travel.usnews.com](http://travel.usnews.com)
3. [bing.com](http://bing.com)
4. [rom.on.ca](http://rom.on.ca)
5. [tripadvisor.com](http://tripadvisor.com)

# Why retrieval-based LMs?

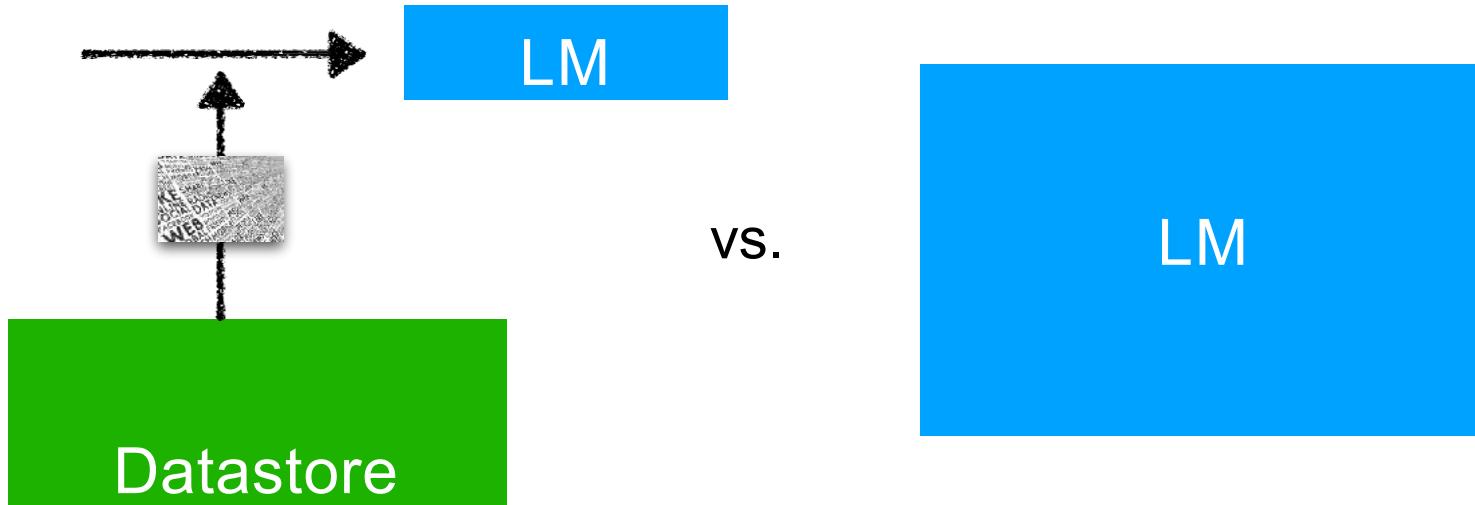


Category	Count
US and international news	109
Log files and error reports	79
License, terms of use, copyright notices	54
Lists of named items (games, countries, etc.)	54
Forum or Wiki entry	53
Valid URLs	50
<b>Named individuals (non-news samples only)</b>	<b>46</b>
Promotional content (products, subscriptions, etc.)	45
High entropy (UUIDs, base64 data)	35
<b>Contact info (address, email, phone, twitter, etc.)</b>	<b>32</b>
Code	31
Configuration files	30
Religious texts	25
Pseudonyms	15
Donald Trump tweets and quotes	12
Web forms (menu items, instructions, etc.)	11
Tech news	11
Lists of numbers (dates, sequences, etc.)	10

Individualization on private data by storing it in the datastore

LLMs are **\*large\*** and expensive to train and run

# Why retrieval-based LMs?



**Long-term goal:** can we possibly reduce the **training** and **inference costs**, and scale down the size of LLMs?

e.g., RETRO (Borgeaud et al., 2021): “obtains comparable performance to GPT-3 on the Pile, despite using **25x fewer parameters**”

# A Retrieval-based LM: Definition

---

A language model (LM) that uses  
**an external datastore at test time**

# Typical LMs



The capital city of Ontario is **Toronto**



LM

# Training time

The capital city of Ontario is \_\_\_\_\_



LM

# Test time / Inference

# Retrieval-based LMs



The capital city of Ontario is **Toronto**



LM

Training time



The capital city of Ontario is \_\_\_\_\_



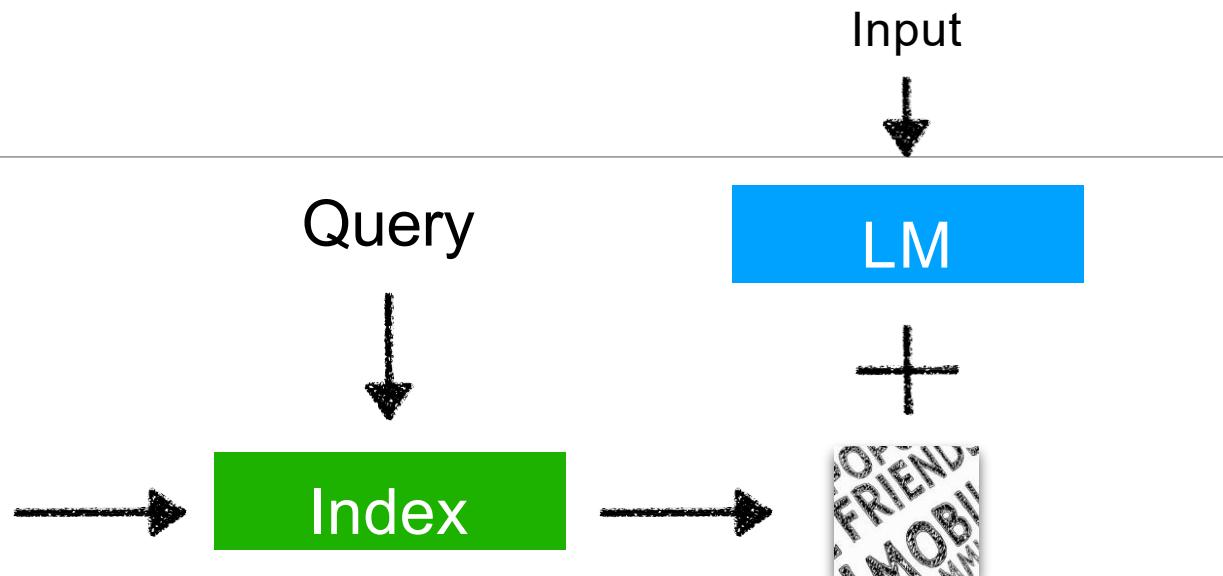
LM

Test time / Inference

# Inference



# Datastore

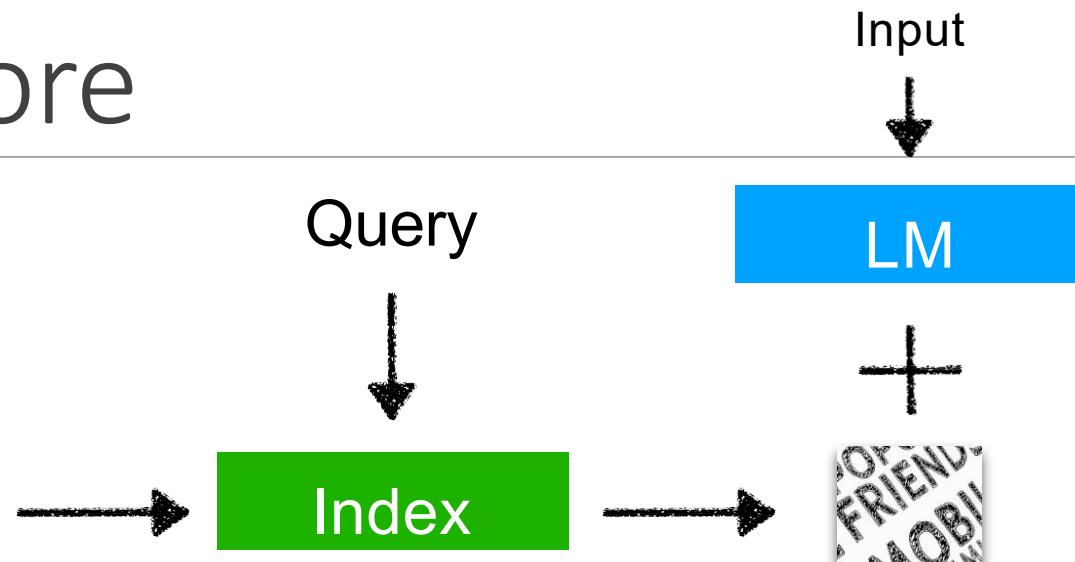


# Inference: Datastore

More recently  
people have used  
structured data

## Datastore

# Raw text corpus (originally)



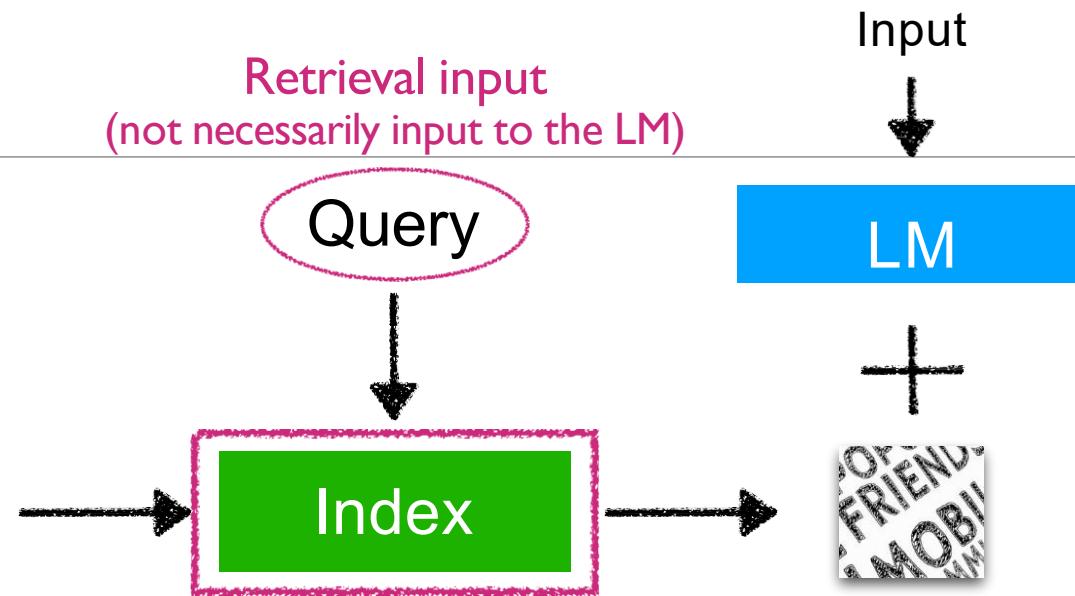
At least billions~trillions of tokens  
Not labeled datasets  
Not structured data (knowledge bases)

# Inference: Index



# Datastore

## Retrieval input (not necessarily input to the LM)



Find a small subset of elements in a datastore  
that are the most similar to the query

# Inference: Index

---

**Goal:** find a small subset of elements in a datastore that are the most similar to the query

**sim:** a similarity score between two pieces of text

# Inference: Index

**Goal:** find a small subset of elements in a datastore that are the most similar to the query

**sim:** a similarity score between two pieces of text

**Example**  $\text{sim}(i, j) = \frac{\text{tf}_{i,j} \times \log \frac{N}{\text{df}_i}}{\# \text{ of occurrences of } i \text{ in } j}$  # of total docs  
# of docs containing

Remember cosine similarity from our discussion of word embeddings

$$\frac{\vec{a} \cdot \vec{b}}{|\vec{a}| |\vec{b}|}$$

**Example**  $\text{sim}(i, j) = \underline{\text{Encoder}(i)} \cdot \underline{\text{Encoder}(j)}$   
Maps the text into an -dimensional vector

# Inference: Index

---

**Goal:** find a small subset of elements in a datastore that are the most similar to the query

**sim:** a similarity score between two pieces of text

Can be a totally separate research area  
on how to do this fast & accurate

---

**Index:** given  $q$ , return  $\arg\max_{d \in \mathcal{D}} \text{sim}(q, d)$  through fast nearest neighbor search

$k$  elements from a datastore

# Software: FAISS, Distributed FAISS, SCaNN, etc...

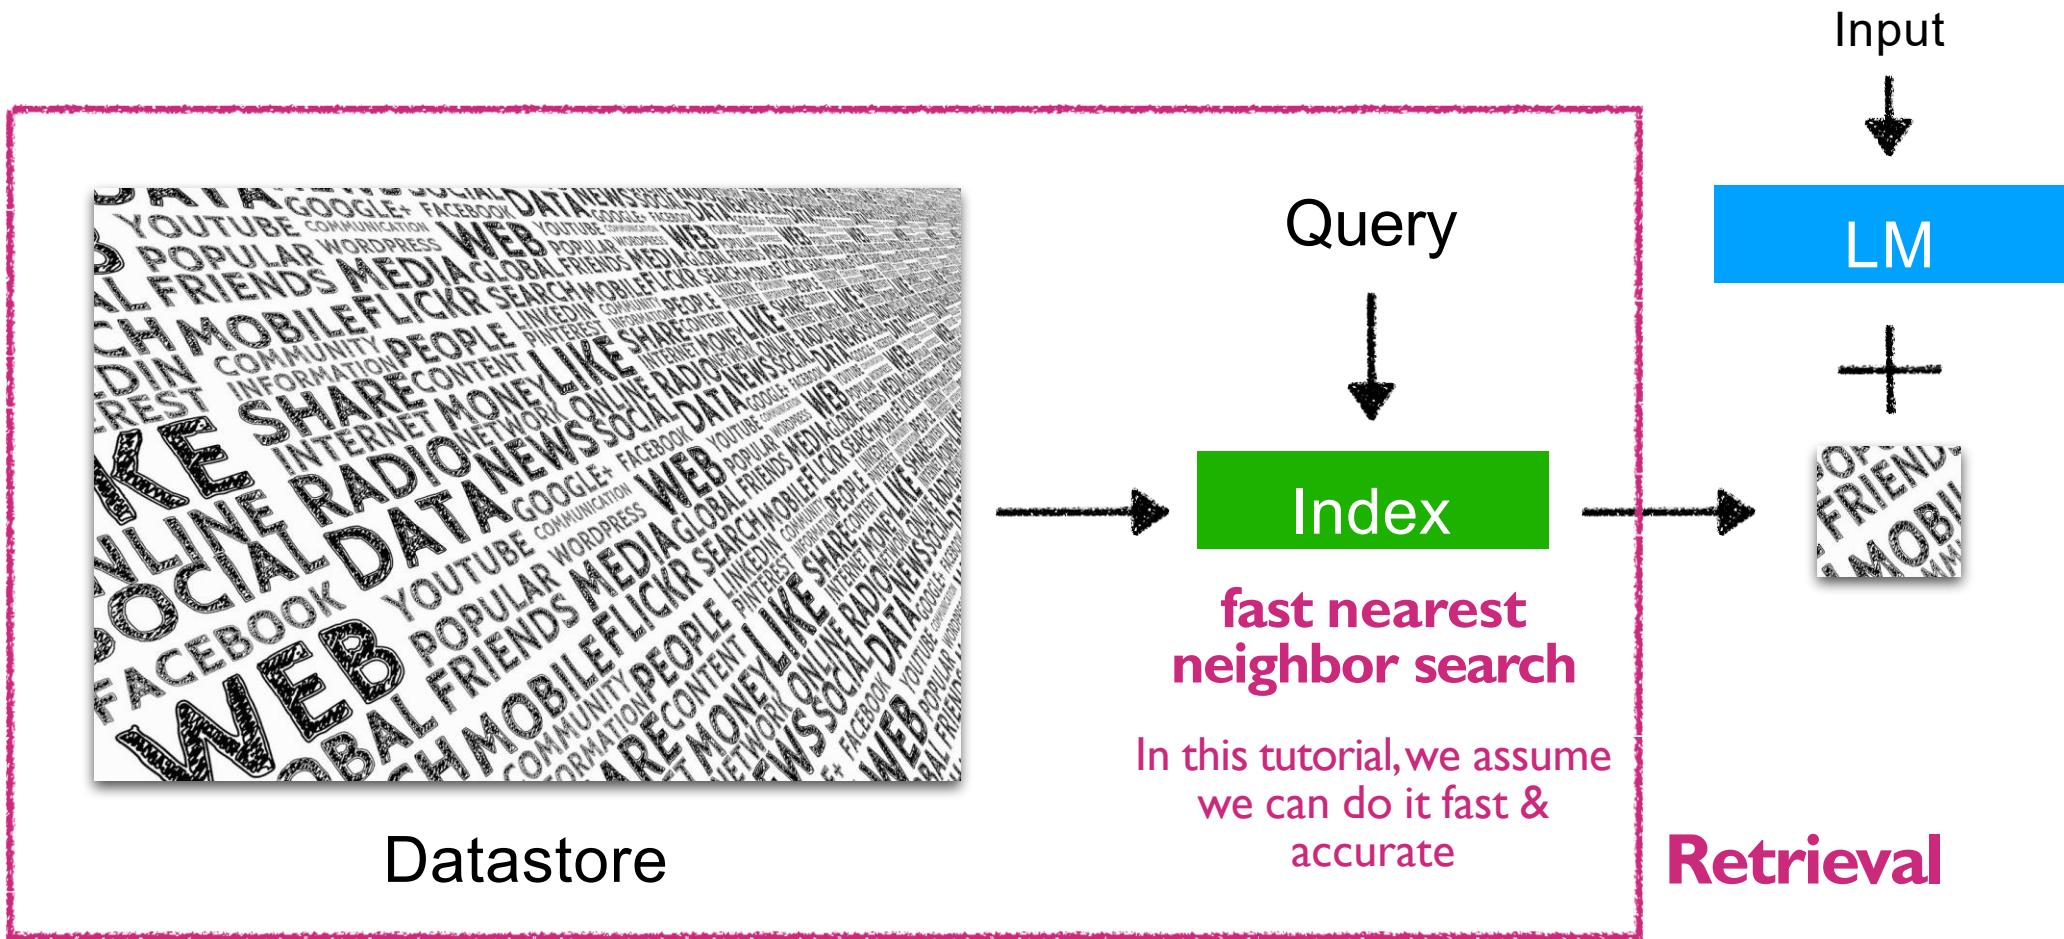
Method	Class name	<code>index_factory</code>	Main parameters	Bytes/vector	Exhaustive	Comments
Exact Search for L2	<code>IndexFlatL2</code>	<code>"Flat"</code>	<code>d</code>	<code>4*d</code>	yes	brute-force
Exact Search for Inner Product	<code>IndexFlatIP</code>	<code>"Flat"</code>	<code>d</code>	<code>4*d</code>	yes	also for cosine (normalize vectors beforehand)
Hierarchical Navigable Small World graph exploration	<code>IndexHNSWFlat</code>	<code>"HNSW,Flat"</code>	<code>d, M</code>	<code>4*d + x * M * 2 * 4</code>	no	
Inverted file with exact post-verification	<code>IndexIVFFlat</code>	<code>"IVFx,Flat"</code>	<code>quantizer, d, nlists, metric</code>	<code>4*d + 8</code>	no	Takes another index to assign vectors to inverted lists. The 8 additional bytes are the vector id that needs to be stored.
Locality-Sensitive Hashing (binary flat index)	<code>IndexLSH</code>	-	<code>d, nbits</code>	<code>ceil(nbites/8)</code>	yes	optimized by using random rotation instead of random projections
Scalar quantizer (SQ) in flat mode	<code>IndexScalarQuantizer</code>	<code>"SQ8"</code>	<code>d</code>	<code>d</code>	yes	4 and 6 bits per component are also implemented.
Product quantizer (PQ) in flat mode	<code>IndexPQ</code>	<code>"PQx", "PQ"m"x"nbits</code>	<code>d, M, nbits</code>	<code>ceil(M * nbits / 8)</code>	yes	
IVF and scalar quantizer	<code>IndexIVFScalarQuantizer</code>	<code>"IVFx,SQ4", "IVFx,SQ8"</code>	<code>quantizer, d, nlists, qtype</code>	<code>SQfp16: 2 * d + 8, SQ8: d + 8 or SQ4: d/2 + 8</code>	no	Same as the <code>IndexScalarQuantizer</code>
IVFADC (coarse quantizer+PQ on residuals)	<code>IndexIVFPQ</code>	<code>"IVFx,PQ"y"x"nbits</code>	<code>quantizer, d, nlists, M, nbits</code>	<code>ceil(M * nbites/8)+8</code>	no	
IVFADC+R (same as IVFADC with re-ranking based on codes)	<code>IndexIVFPQR</code>	<code>"IVFx,PQy+z"</code>	<code>quantizer, d, nlists, M, nbits, M_refine, nbits_refine</code>	<code>M+M_refine+8</code>	no	

Exact Search

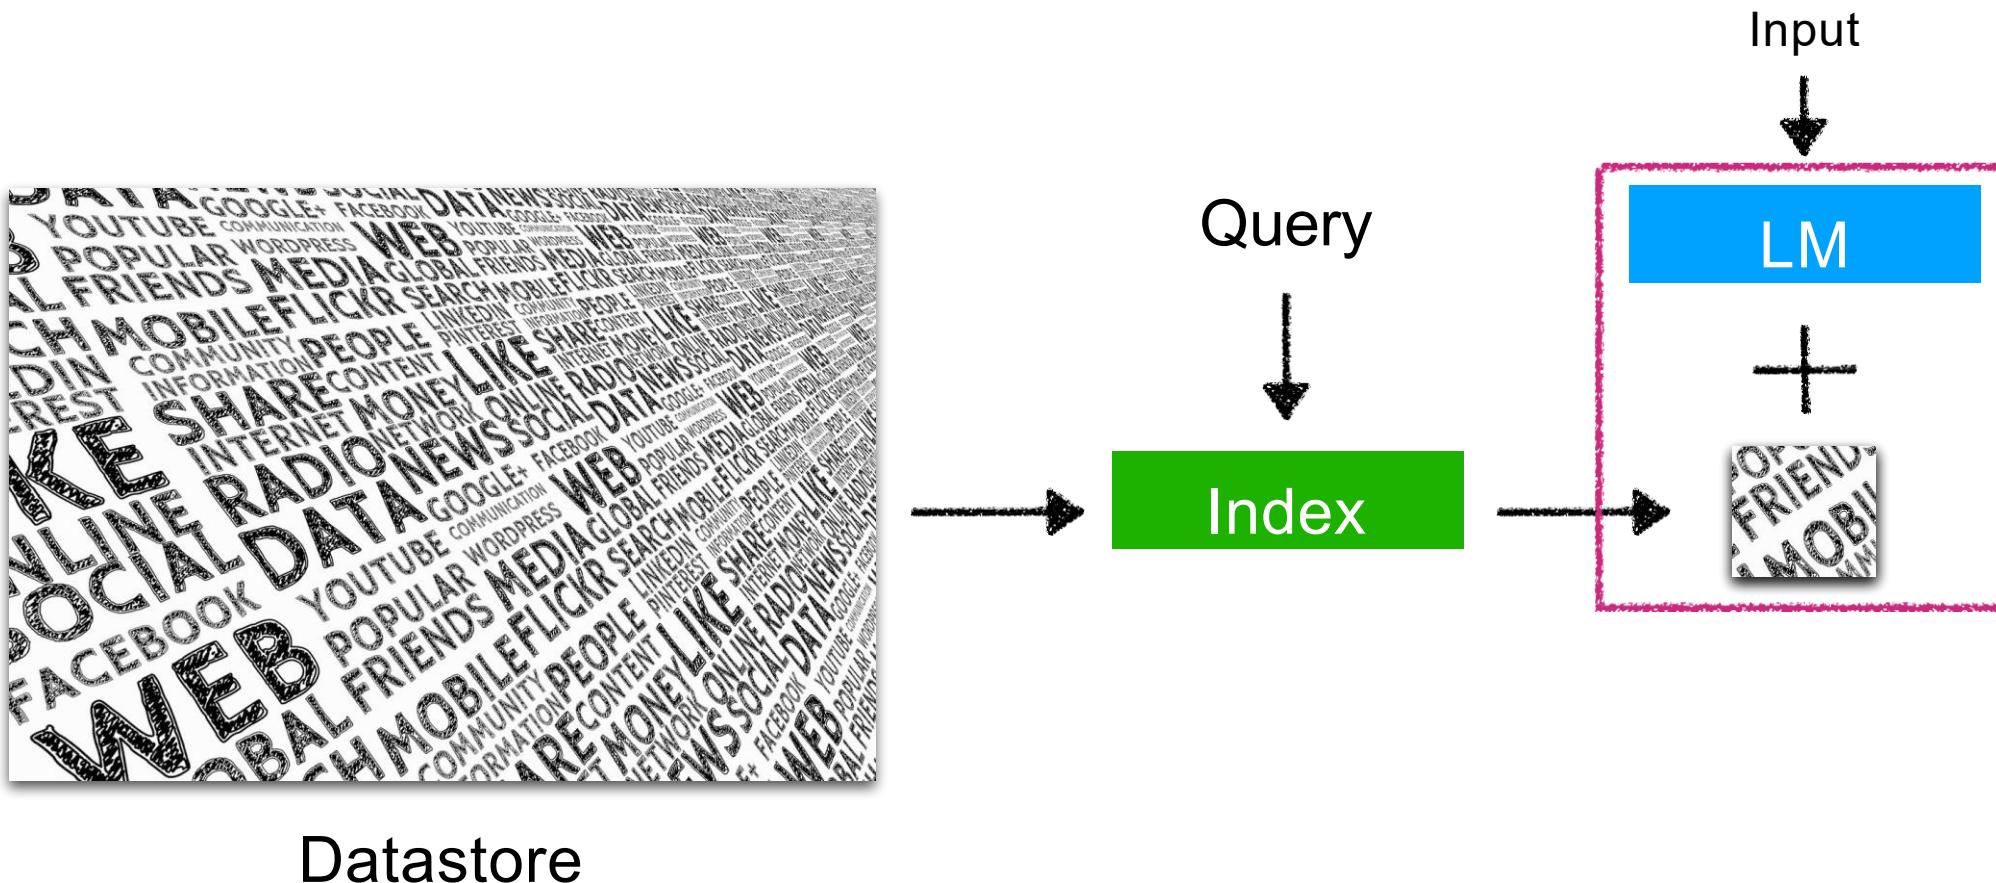
Approximate Search  
(Relatively easy to scale to ~1B elements)

More info: <https://github.com/facebookresearch/faiss/wiki>

# Inference: Search



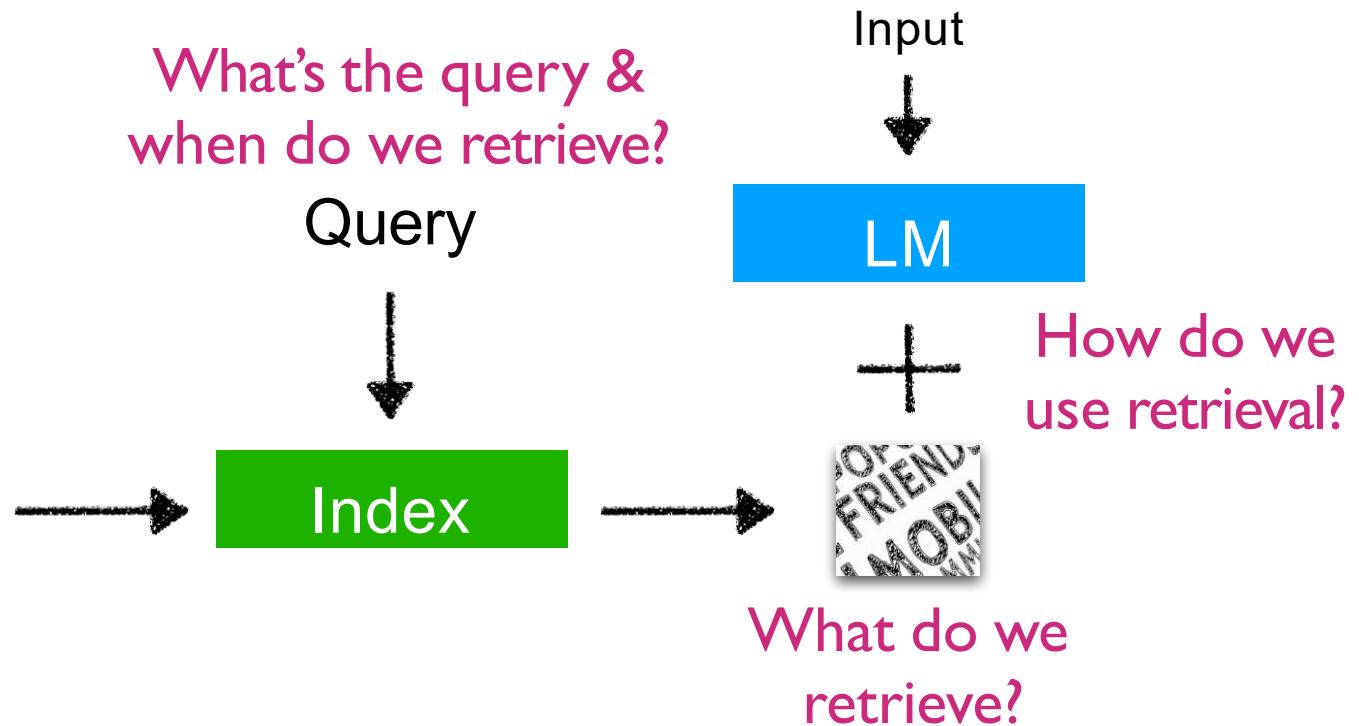
# Inference: Search



# Variations of RAG



# Datastore



# Variations of RAG

**What** to retrieve?

Query



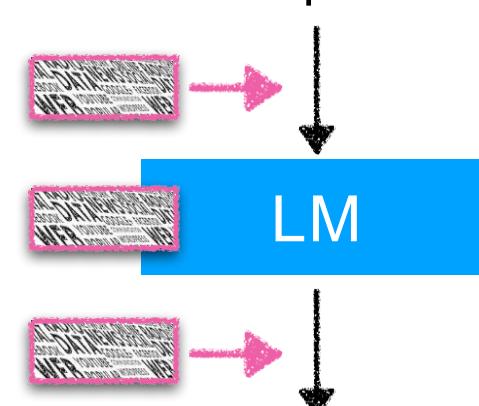
Text chunks (passages)?

Tokens?

Something else?

**How** to use retrieval?

Input



**When** to retrieve?

w/ retrieval

The capital city of Ontario is Toronto.

w/ retrieval w/ r w/r w/r w/r w/r

The capital city of Ontario is Toronto.

w/ retrieval

The capital city of Ontario is Toronto.

# In-Class Activity

Skim the paper assigned to you

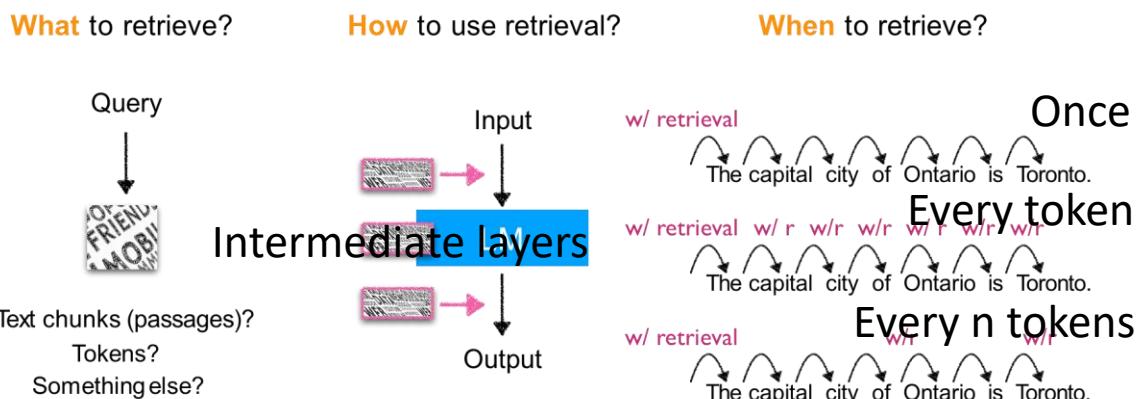
In your paper, find the answers to these questions

**What to retrieve?**

**How to use retrieval?**

**When to retrieve?**

Share what you learned with your table



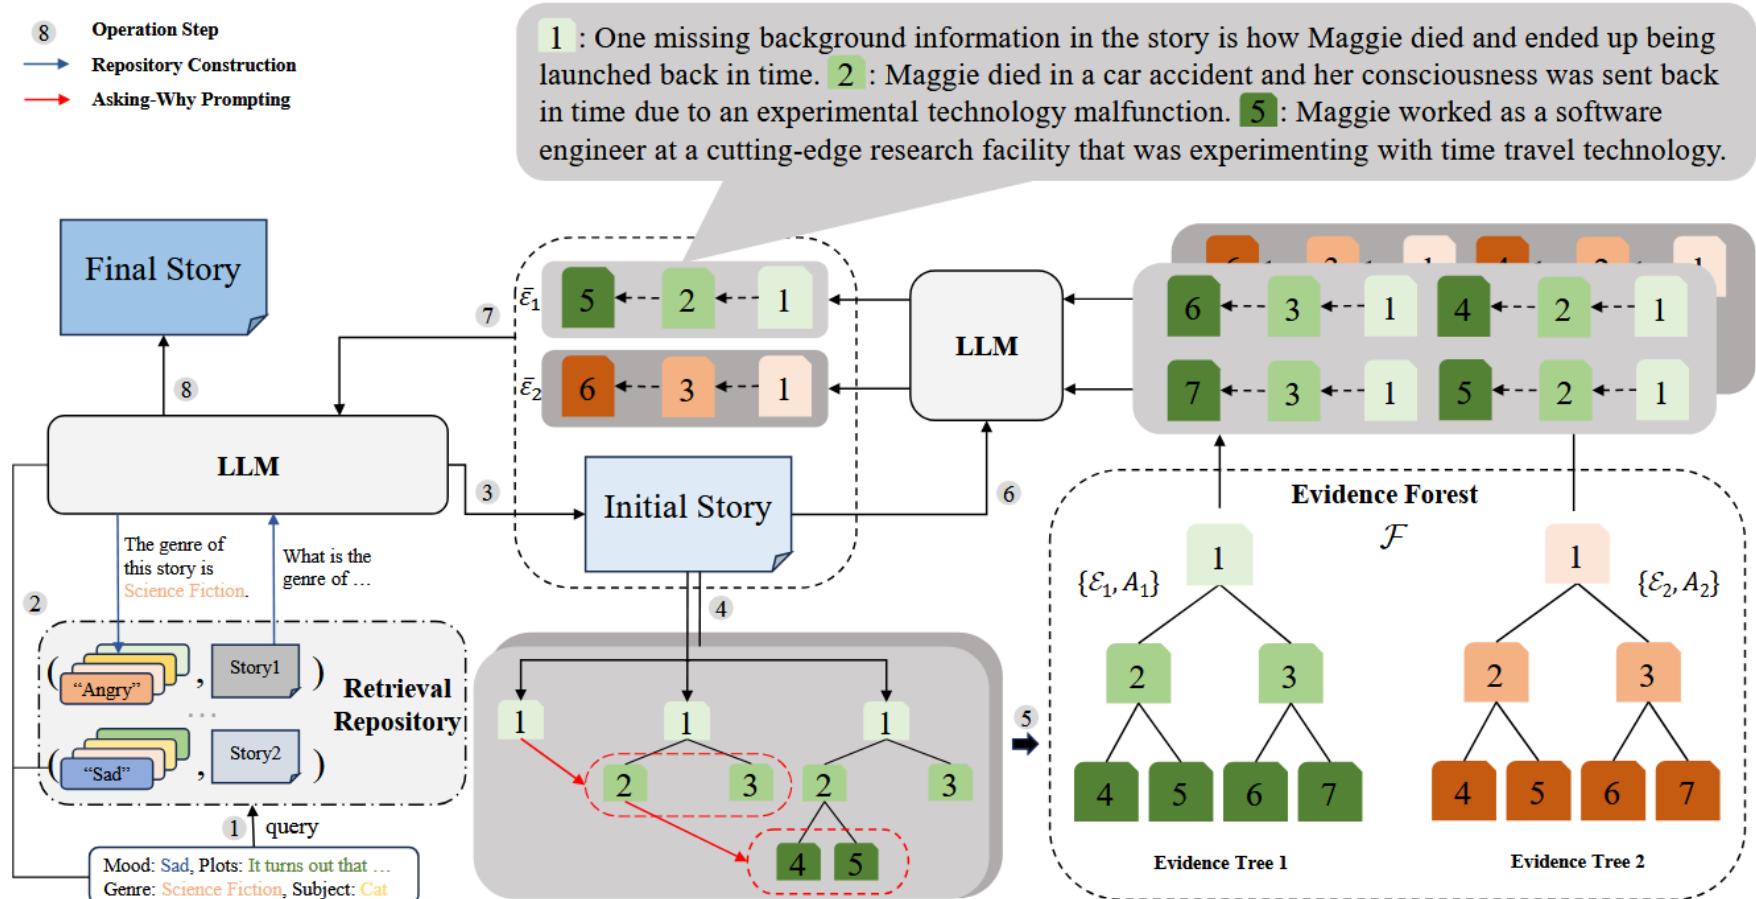
# Say Anything (2009)

<i>User written sentence</i>	<i>Matching sentence</i>	<i>Computer generated sentence</i>
It all started easy enough.	It all started out normal enough.	Got up, had lunch, got ready went to work etc.
But then I noticed something very strange.	Then, as I flipped through the test, I realized something very strange.	Mr. Sipiera gave me points on the blank page.
Why would he do that? What was he up to?	Why would he do that?	Because he was planning to kill Sonya for her knowledge of the kidnapping plot.
I decided to get to Sonya first and run away with her to a Greek island.	We left pretty much after this and took a ferry back to Hong Kong island.	Our original plan to go to Lamma island was ditched, as we had inadvertently spent half the day on Lantau island.
But we settled for Corfu, which has a beautiful little hotel on the edge of the old town.	But we settled for the simpler name.	We had dogs (and cats) when I was growing up but my wife and sons have never had a dog before.
Corfu was full of cats and dogs, and I ended up having to look at my feet a lot when walking around town.	and i ended up having to save the day.	It's too long to list the details of the whole operation.
Sonya would take on the secret identity of a scuba instructor, and I'd be a fisherman.	"Before I even got so far as having a proper breakdown in realizing that fantasy is not reality, I found that I had the secret identity of a beetle."	"But that could have been very useful, right?
It meant we could spend a lot of secret time together on the jetty beyond the coastline.	She does spend a lot of time at the bar where they all hang out.	She has told me some of the men she hangs with have come on to her.
I told her I'd kill them just like that old math teacher.	I told her I'd come help her clean and whatever else she needs.	Once she's gone, I'm pretty much on my own.
Fishing all day, looking out for the police.	This is bad for the police.	If they stopped to survey the situation they could only see what was there -- none of the men were armed.
It was a great secret life.		

Before neural methods,  
retrieval was used for  
generating stories

# GROVE (2023)

- ⑧ Operation Step
- Repository Construction
- Asking-Why Prompting



# BERALL (2024) (to be presented on Tuesday)

---

