

Planning

Lara J. Martin (she/they)

<https://laramartin.net/interactive-fiction-class>

Slides adapted from Chris Callison-Burch

Learning Objectives

Identify the components of a planning problem and how they are used in PDDL

Distinguish between search and planning (and MDPs)

Discover limitations and benefits of planning for storytelling

Assess how LLMs and planning can work together for storytelling

Review:

Formal Definition of a Search Problem

1. **States:** a set S

2. An **initial state** $s_i \in S$

3. **Actions:** a set A

$\forall s$ **Actions**(s) = the set of actions that can be executed in s .

4. **Transition Model:** $\forall s \forall a \in \text{Actions}(s)$

Result(s, a) $\rightarrow s_r$

s_r is called a **successor** of s

$\{s_i\} \cup \text{Successors}(s_i)^* = \text{state space}$

5. **Path cost** (Performance Measure):

Must be additive, e.g. sum of distances, number of actions executed, ...

$c(x, a, y)$ is the step cost, assumed ≥ 0

◦ (where action a goes from state x to state y)

6. **Goal test:** **Goal**(s)

s is a goal state if **Goal**(s) is true.

Can be implicit, e.g. **checkmate**(s)

Planning

Planning: The process of **searching** for a plan

- This is why we can use algorithms like BFS to find plans
- “Plain” state-based search is useful when we just want to get to the goal (efficiently); planning is useful when we care about the *path*
- What we think of as “planning” is a combination of search and logic

Plan: The result of planning; a sequence of steps from the initial state to a goal state


Policy: A collection of transition functions ($\text{Result}(s, a) \rightarrow s_r$) that tell the agent what action it should take for a given state

- This will become more relevant when talking about reinforcement learning

Classical Planning

Classical planning: The task of finding a sequence of action to accomplish a goal in an environment that:

- Is deterministic
- Is fully observable
- Contains a single agent
- Has a single initial state
- Is discrete



The solution to any problem in such an environment is a **fixed sequence of actions**.

More complicated planning

In environments that are

- Nondeterministic
- Partially observable
- Etc.

Nondeterministic actions
(with assigned probabilities)
turn classical planning
problems into an MDP!

The solution must recommend different future actions depending on the what percepts it receives. This could be in the form of a branching strategy.



Planning languages

PDDL (Planning Domain Definition Language)

STRIPS (Stanford Research Institute Problem Solver)

ADL (Action Description Language)

...

We'll focus on PDDL

PDDL breaks the planning problem into a **domain** and a **problem description**

The **domain** is consistent across problems (e.g., the description of the environment)

The **problem** defines what is going to be planned over

Domain

```
(define (domain action-castle)
  (:requirements :strips :typing)
  (:types
```

Domain name

```
    player location direction
    fish crown - item
  )
```

Object Types (can be hierarchical)

```
(:action go
```

Parameters
(variables)

```
  :parameters (?dir - direction ?p - player ?l1 - location ?l2 - location)
```

Preconditions

```
  :precondition (and (at ?p ?l1)
                     (connected ?l1 ?dir ?l2)
                     (not (blocked ?l1 ?dir ?l2)))
```

Effects

```
  :effect (and (at ?p ?l2)
               (not (at ?p ?l1)))
```

Logical statements

```
)
```

```
(:action get
```

```
  :parameters (?item - item ?p - player ?l1 - location )
```

```
  :precondition (and (at ?p ?l1)
                     (at ?item ?l1))
```

```
  :effect (and (inventory ?p ?item)
               (not (at ?item ?l1)))
```

```
))
```

Actions

Problem

Objects
(the atoms)

```
(define (problem navigate-to-location)
  (:domain action-castle)

  (:objects
    npc - player
    cottage gardenpath fishingpond gardenpath
    windingpath talltree drawbridge courtyard
    towerstairs tower dungeonstairs dungeon
    greatfeastinghall throneroom - location
    in out north south east west up down - direction
  )

  (:init
    (at npc cottage)
    (connected cottage out gardenpath)
    (connected gardenpath in cottage)
    (connected gardenpath south fishingpond)
    (connected fishingpond north gardenpath)
  )

  (:goal (and (at npc throneroom) (sitting npc throne)))
)
```

Problem name

What domain to
use

Initial State

Goal

State Representation

In PDDL, a **state** is represented as a **conjunction** of **logical sentences** that are **ground atomic fluents**.

Ground means
they contain no
variables

Atomic sentences
logical statements
that can't be
broken down

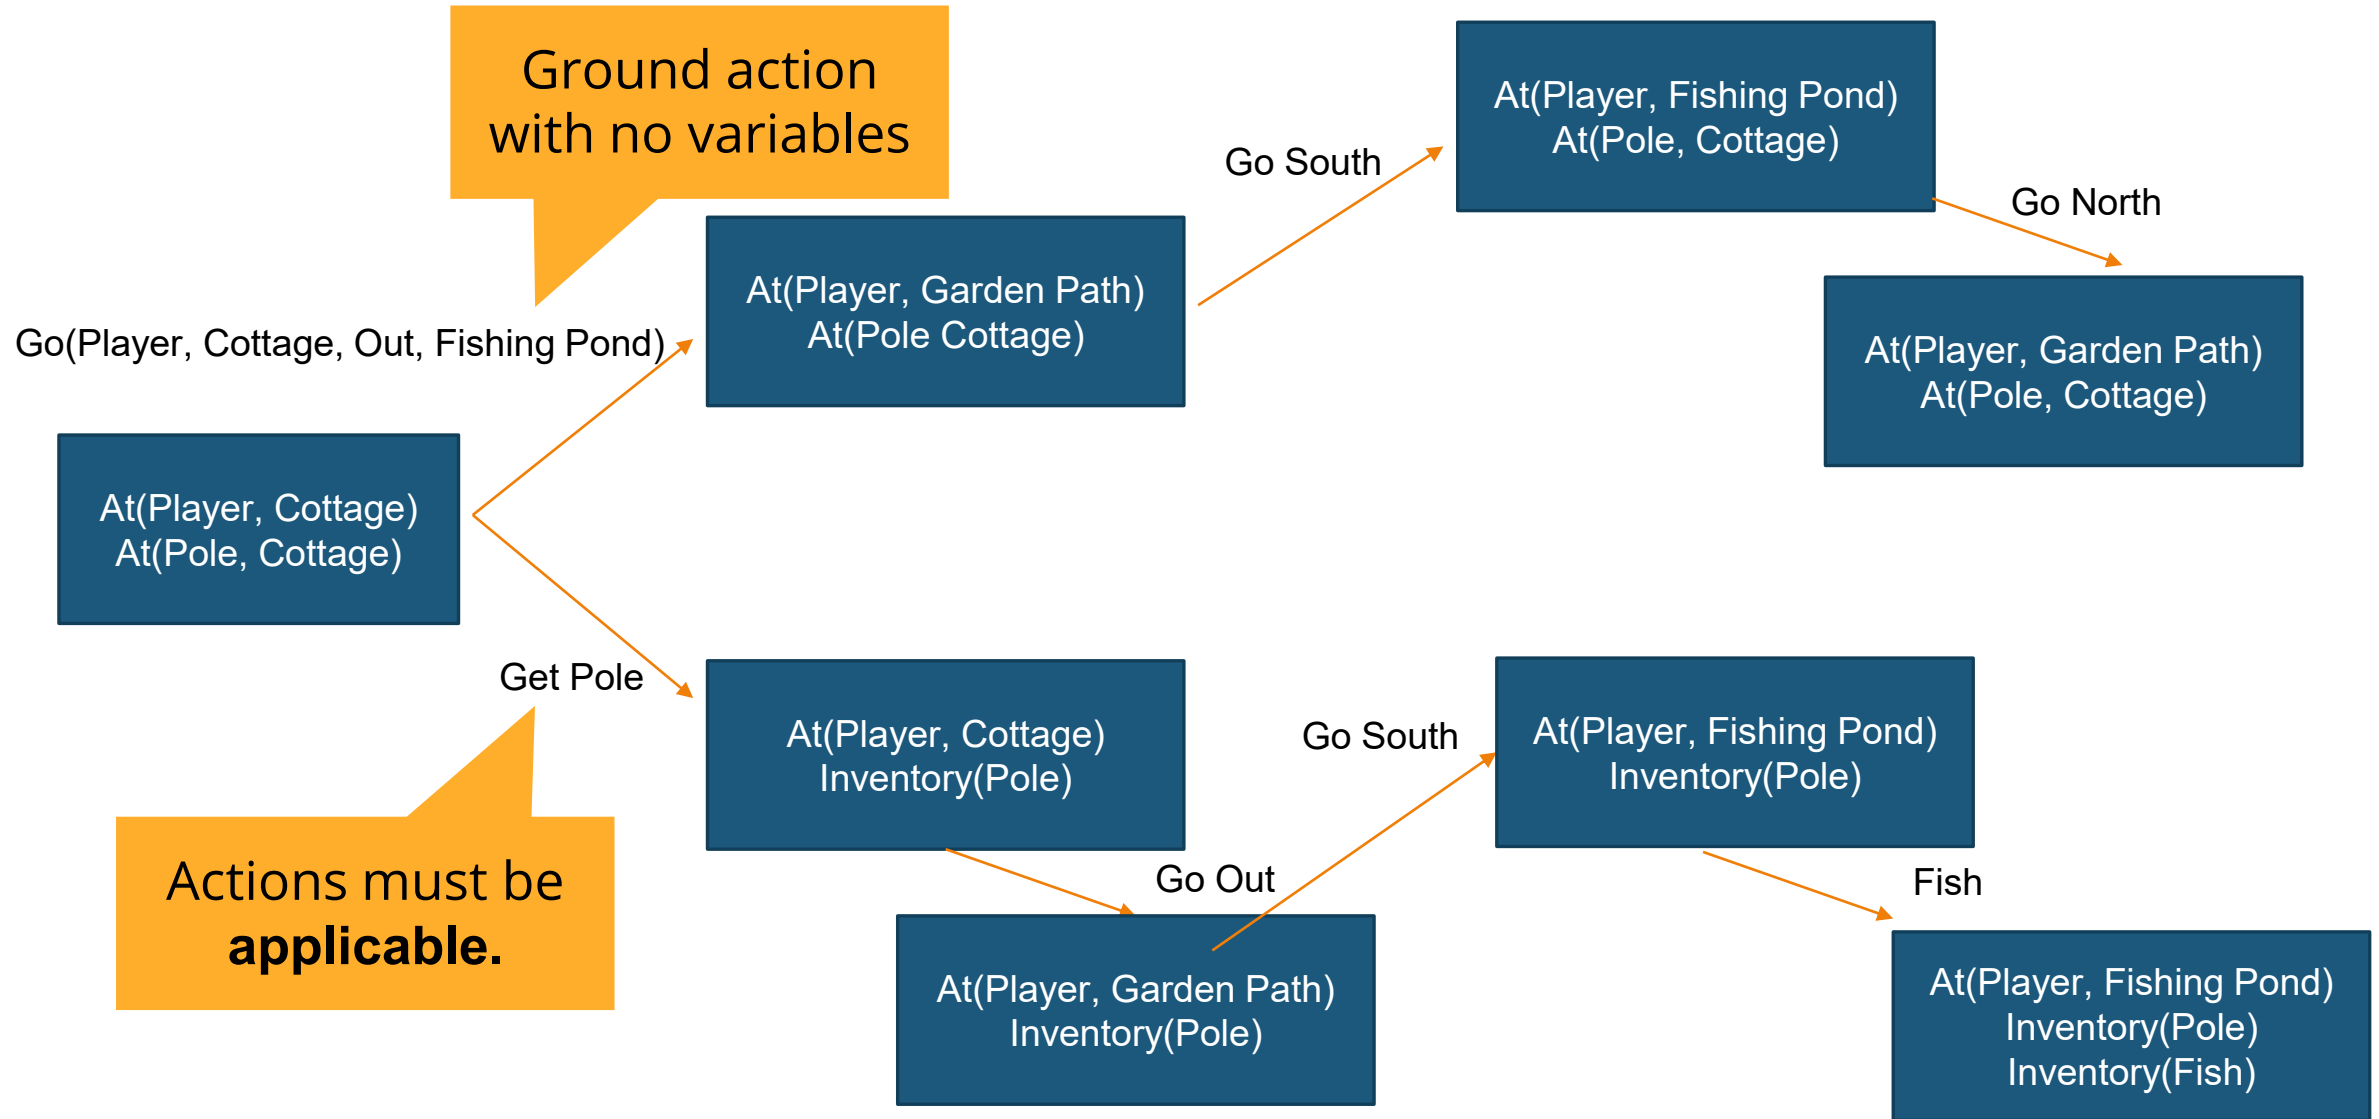
Fluent means an aspect
of the world that can
change over time

E.g.,


(*connected* gardenpath in cottage)

We make the *closed world assumption*: any fluent not mentioned is false

Forward State-Space Search for Planning



Backward State-Space Search for Planning (aka Regression Search)



Given a goal **g** and action **a**, the **regression** from g to a gives a state **g'** description whose literals are given by:

$$\text{POS}(g') = (\text{POS}(g) - \text{ADD}(a)) \cup \text{POS}(\text{Preconditions}(a))$$
$$\text{NEG}(g') = (\text{NEG}(g) - \text{DEL}(a)) \cup \text{NEG}(\text{Preconditions}(a))$$

Negative literals in the effects are kept in a **delete list DEL**

Positive literals in the effects are kept in an ADD list

Backward State-Space Search for Planning (aka Regression Search)

Given a goal **g** and action **a**, the **regression** from g to a gives a state **g'** description whose literals are given by:

$$\text{POS}(g') = (\text{POS}(g) - \text{ADD}(a)) \cup \text{POS}(\text{Preconditions}(a))$$

$$\text{NEG}(g') = (\text{NEG}(g) - \text{DEL}(a)) \cup \text{NEG}(\text{Preconditions}(a))$$

Or simply:

$$g' = (g - \text{effects}(a)) \cup \text{Preconditions}(a)$$

Partial-Order Planning

Keep a **partial order of steps** and only commit to an ordering when forced to

For example:

- Go north
- Pick up sword; Pick up lantern
- Go west

Review:

Partial Order Causal Link (POCL) planning

Figure 1: Example CPOCL Problem and Domain

Initial: $\text{single}(A) \wedge \text{single}(B) \wedge \text{single}(C) \wedge$
 $\text{loves}(A, C) \wedge \text{intends}(A, \text{married}(A, C)) \wedge$
 $\text{loves}(B, C) \wedge \text{intends}(B, \text{married}(B, C)) \wedge \text{has}(B, R)$

Goal: $\text{married}(A, C)$

lose(?p, ?i)
 A: \emptyset
 P: $\text{has}(\text{?p}, \text{?i})$
 E: $\text{lost}(\text{?i}) \wedge \neg \text{has}(\text{?p}, \text{?i})$

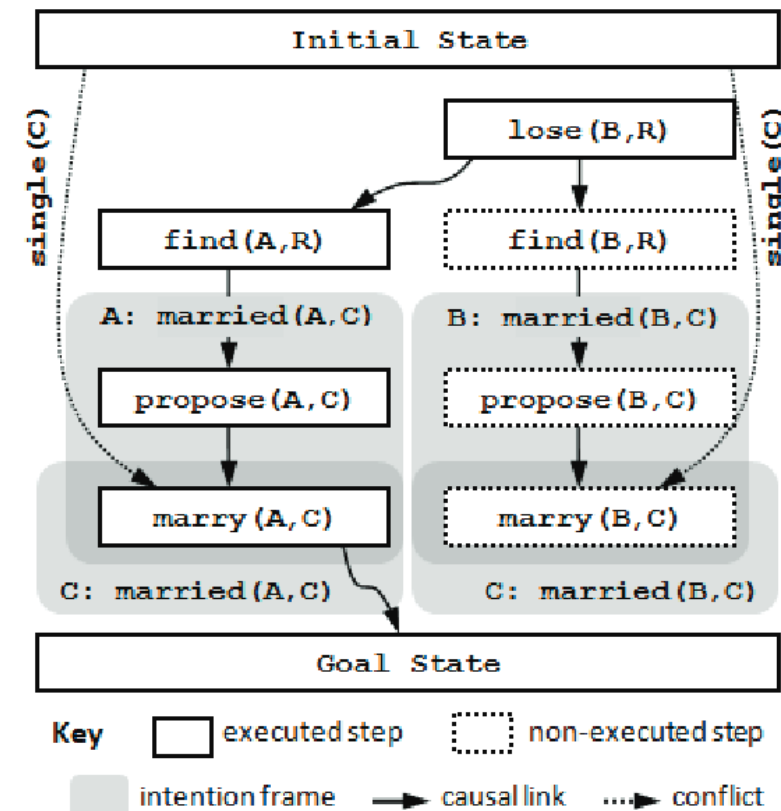
give(?p1, ?p2, ?i)
 A: ?p1 ?p2
 P: $\text{has}(\text{?p1}, \text{?i})$
 E: $\text{has}(\text{?p2}, \text{?i}) \wedge \neg \text{has}(\text{?p1}, \text{?i})$

find(?p, ?i)
 A: \emptyset
 P: $\text{lost}(\text{?i})$
 E: $\text{has}(\text{?p}, \text{?i}) \wedge \neg \text{lost}(\text{?i})$

marry(?b, ?g)
 A: ?b ?g
 P: $\text{loves}(\text{?b}, \text{?g}) \wedge \text{loves}(\text{?g}, \text{?b}) \wedge \text{single}(\text{?b}) \wedge \text{single}(\text{?g})$
 E: $\text{married}(\text{?b}, \text{?g}) \wedge \neg \text{single}(\text{?b}) \wedge \neg \text{single}(\text{?g})$

propose(?b, ?g)
 A: ?b
 P: $\text{loves}(\text{?b}, \text{?g}) \wedge \text{has}(\text{?b}, R)$
 E: $\text{loves}(\text{?g}, \text{?b}) \wedge \text{intends}(\text{?g}, \text{married}(\text{?b}, \text{?g}))$

Figure 2: Example CPOCL Plan



Think-Pair-Share

What are some limitations of planning for storytelling?

What are some benefits?


Neural Planning

Types of Neural Planning

Generation of planning language code

- To be run through planner

PDDLGO



Iterative creation of PDDL problem

Treat as partially-observed

Keep regenerating until plan succeeds

Types of Neural Planning

Generation of planning language code

- To be run through planner

Generation of planning-esque components

Neural Story Planning

Algorithm 1: Neural Plot Planner

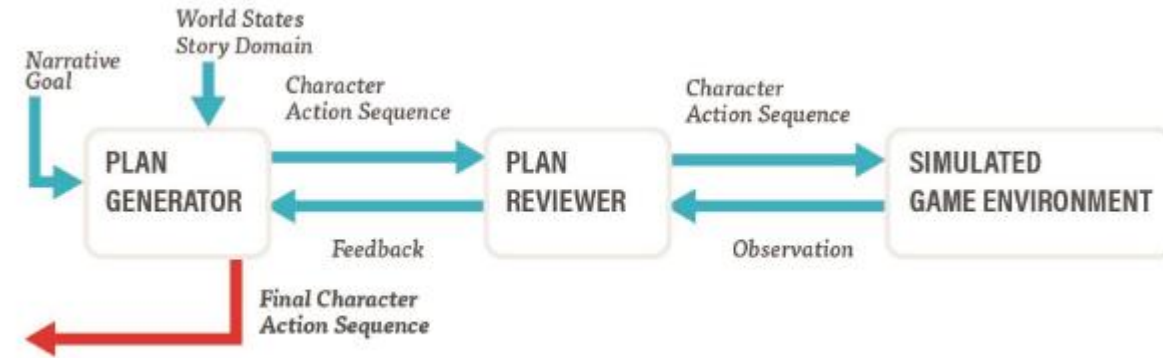
```
1: Input: ending event sentence  $g$ ; Initial conditions  $I$ .
2: Initialize a plan  $P \leftarrow \emptyset$ ; Initialize  $queue \leftarrow \{g\}$ .
3: while  $queue \neq \emptyset$  do
4:   Let  $event \leftarrow \text{pop}(queue)$ 
5:   Let  $context \leftarrow$  sequence of events collected by running a breadth-first
     search from  $event$  to  $g$ .
6:   Let  $\Lambda \leftarrow$  all satisfied preconditions
7:   Let  $\Gamma \leftarrow \text{generate preconds}(event)$  for each character in  $event$ 
8:   if adding any precondition in  $\Gamma$  creates a cycle then
9:     remove  $event$  from  $P$ .
10:     $\Gamma \leftarrow$  unsatisfied preconditions due to removing  $event$ .
11:   for each  $c \in \Gamma$  do
12:     if  $c \in I$  or  $c$  meets conditions for not being expanded then
13:        $P \leftarrow P \cup \{nil \xrightarrow{c} event\}$   $\triangleright$  Dangling precondition
14:     else if there exists a precondition  $c' \in \Lambda$  that is similar to  $c$  then
15:        $event' \leftarrow$  event that satisfies  $c'$ 
16:        $P \leftarrow P \cup \{event' \xrightarrow{c'} event\}$   $\triangleright$  Reuse precondition
17:     else
18:        $event' \leftarrow \text{generate event}(c, context)$ 
19:        $P \leftarrow P \cup \{event' \xrightarrow{c} event\}$   $\triangleright$  Satisfy with new event
20:    $queue \leftarrow queue \cup \{event'\}$ 
```

Generates preconditions for characters

- Items needed
- Locations
- Item states
- Context (How)
- Interactions with others
- Reasons

Partial-order-planning–inspired generation

StoryVerse



STORY DOMAIN

CHARACTERS



LOCATIONS



ACTION SCHEMA

MoveTo(X)
SlipIntoWater()
DrownToDeath()
Save(X)
TryToKill(X)
Kill(X)
Think(X)

CHARACTER SIMULATION

CHARACTER SIMULATOR INITIALIZES THE STATE

I should head to the **bank** to find some food and maybe meet other creatures. It's always nice to socialize and share the latest news.

MoveTo(Bank)

I feel like visiting the **bank** today. It's always nice to meet new friends, and maybe I'll see the **ant** there.

MoveTo(Bank)

I should head to the **bank**. It's a good spot to find targets. I might find some animals there. Being close to water could lead to more opportunities.

MoveTo(Bank)

INSTANTIATED ACTS

ACT 1

SOME CHARACTER GOT INTO AN ACCIDENT

I see some delicious leaves across the **brook**. I should go there to eat.

MoveTo(Brook)

Oh no! The edge is very slippery! I'm falling into the water!

SlipIntoWater()

ACT 2

SOME CHARACTER SAVED [ANT]

I should go to the **brook**. Maybe I can find the **ant** there and ensure its safe.

MoveTo(Brook)

I see the **ant** struggling in the water. I must save it before it's too late.

Save(Ant)

ACT 3

[DOVE] GOT INTO A DIFFERENT ACCIDENT

I should move closer to the **brook**. I might get a better shot at my target there.

MoveTo(Brook)

There's the **dove**! Saving the **ant**. It's distracted. Now it's my chance.

TryToKill(Dove)

ACT 4

[ANT] SAVED [DOVE]

I see the **hunter** aiming at the **dove**. I must do something quickly.

Think(...)

If I bite the **hunter**, maybe I can distract him and save the **dove**.

TryToKill(Dove)

CHARACTER SIMULATION

THE END STATE IS REACHED, SIMULATION CONTINUES

Now that I've had a bit of an adventure and was saved by the **dove**, I should go to the **oak tree** to find some food. It's safer there.

MoveTo(OakTree)

Now that the ant is safe, I should check on the **oak tree**. It's a good place to rest and maybe I'll find more friends there.

MoveTo(OakTree)

The **oak tree** next to the **brook** could be a promising spot to find game. Plus, there's a **dove** there. Maybe I can learn more and find an opportunity.

MoveTo(OakTree)

Figure 2: Two example story domains - The Ville (top) and Ant & Dove (bottom) - together with instantiated versions of the four abstract acts from Figure 1. Note that the text for narrations, dialogs, and monologues is all generated by LLMs.

Types of Neural Planning


Generation of planning language code

- To be run through planner

Generation of planning-esque components^{†c}

Use of LLM as *informal* planner

- Like the procedures work we looked at



Types of Neural Planning

Generation of planning language code

- To be run through planner

Generation of planning-esque components

Use of LLM as *informal* planner

- Like the procedures work we looked at
- Or using techniques like CoT

```
{ "role": "system", "content": "# You are a helpful fiction writer assistant." },
{ "role": "user", "content": f"{plot}\n
Summarize the plot above into a plot tree of
{'at most 6' if num_nodes == '' else num_nodes}
nodes with each node containing the state and goal of {char_name}, and the key
decision that propels the story forward. Each edge should contain a list of
events that lead {char_name} to the state of next node. Also, Given the same
state and goal of {char_name}, imagine an alternate decision that would have led
{char_name} to a different storyline. Output in JSON format with schema:
{JSON_SCHEMA}. Make sure that all important plot points are included in
'edgeEvents' but not in 'state'"
}
```

Table 5: Prompt for generating a tree from the plot (plot-to-tree).

Types of Neural Planning

Generation of planning language code

- To be run through planner

Generation of planning-esque components

Use of LLM as *informal* planner

- Like the procedures work we looked at
- Or using techniques like CoT
- Or guided/hierarchical generation like Plan & Write

Dynamic	Storyline	needed → money → computer → bought → happy
	Story	John <u>needed</u> a computer for his birthday. He worked hard to earn <u>money</u> . John was able to buy his <u>computer</u> . He went to the store and <u>bought</u> a computer. John was <u>happy</u> with his new computer.
Static	Storyline	computer → slow → work → day → buy
	Story	I have an old <u>computer</u> . It was very <u>slow</u> . I tried to <u>work</u> on it but it wouldn't work. One <u>day</u> , I decided to buy a new one. I <u>bought</u> a new computer .