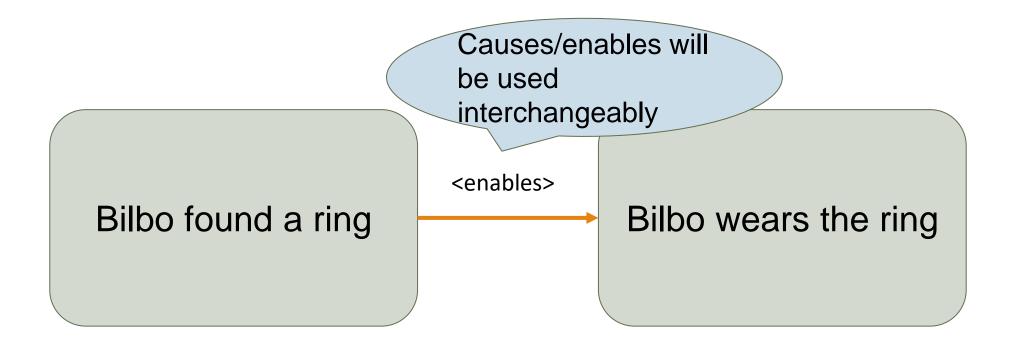
Schedule

- Module 2 Review
- 1 paper presentation
- Begin lecture on search/planning

Module 2 Review

10/8/2024 CMSC 491/691 - INTERACTIVE FICTION AND TEXT GENERATION DR. LARA J. MARTIN

Review: Causal Links



Review: Script

"A standard event sequence" that

- Lays out different paths/options
- Consists of causal chains
- Can be used to leave out tedious details the reader is expected to know
- Can be considered a literary trope or a common social scenario

Review: Principle of Minimal Departure

"This law—to which I shall refer as the principle of minimal departure—states that we reconstrue the central world of a textual universe in the same way we reconstrue the alternate possible worlds of nonfactual statement: as conforming as far as possible to our representation of [the actual world]"

In other words:

The story world is expected to be like the real world, unless otherwise specified

Ryan, M.-L. (1991). Chapter 3: Reconstructing the Textual Universe: The Principle of Minimal Departure. In *Possible Worlds, Artificial Intelligence, and Narrative Theory* (pp. 48–60). Indiana Univ. Press.

Review: Linking Events

PROBABILISTIC

CAUSAL

Occur frequently together (not necessarily because they had to)

Occur because of one another

Example:

I pour dog food in my dog's bowl.

I pet my dog.

Example:

MODULE 2 REVIEW

I pour dog food in my dog's bowl.

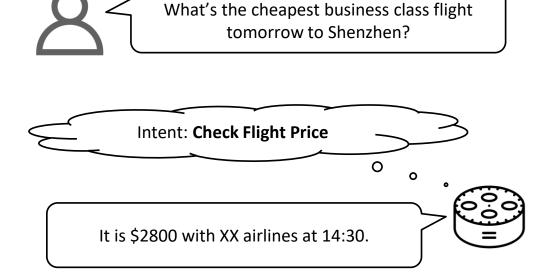
My dog eats dog food.

Review: What are procedures?

- A procedure is "a series of actions conducted in a certain order or manner," as
 defined by Oxford
- A more refined definition: "a series of steps happening to achieve some goal^[1]"
 - Why?
- Examples of procedures: instructions (recipes, manuals, navigation info, how-to guide), algorithm, scientific processes, etc.
 - We focus on instructions, which is human-centered and task-oriented
- Examples of non-procedures: news articles, novels, descriptions, etc.
 - Those are often narrative: events do not have a specific goal
 - The umbrella term is script^[2]

Review: Intent Detection

- Task-oriented dialog systems needs to match an utterance to an intent, before making informed responses
- Sentence classification task
 - Given an utterance, and some candidate intents
 - Choose the correct intent
 - Evaluated by accuracy



Example from Snips (Coucke et al., 2018)

10/8/2024

Utterance: "Find the schedule at Star Theatres."

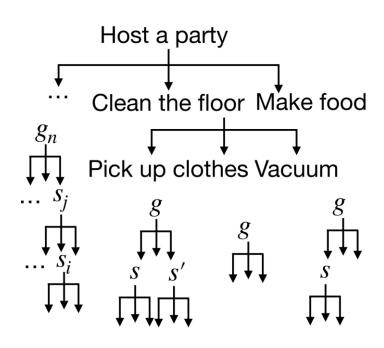
Candidate intents: Add to Playlist, Rate Book, Book Restaurant,

Get Weather, Play Music, Search Creative Work, Search Screening Event

8

Review: Procedures are Hierarchical

- An event can simultaneously be a goal of one procedure, and a step in another
- A procedural hierarchy... So what?
 - Can "explain in more details" by expansion
 - Can shed light on event granularity (why?)
- How do you build such hierarchy?
 - To "host a party", I need to "clean the floor"; to "clean the floor", I need to do what?



Review: Plan-and-Write

<u>Carrie</u> had just learned how to ride a bike. She didn't have a <u>bike</u> of her own. Carrie would <u>sneak</u> rides on her sister's bike. She got <u>nervous</u> on a hill and crashed into a wall. The bike frame bent and Carrie got a deep gash on her <u>leg</u>.

Carrie→bike→sneak→nervous→leg

10/8/2024

Review: Guided Open Story Generation Using Probabilistic Graphical Models

Use discourse representation structure (DRS) parser to get semantic relationships

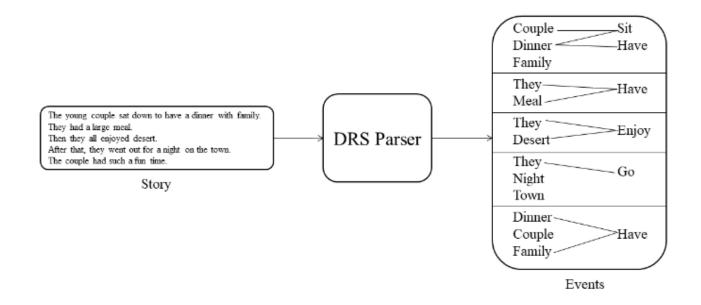


Figure 1: Event Representation using DRS Parser. Note that the words without edges are removed while forming the graphs.

Review: Example of a Probabilistic Event Representation

From sentence, extract event representation:

(subject, verb, direct object, modifier, preposition)

Original sentence: yoda uses the force to take apart the platform

Events:

```
yoda use force \emptyset \emptyset yoda take_apart platform \emptyset \emptyset
```

Generalized Events:

```
<PERSON>0 fit-54.3 power.n.01 Ø Ø
```

<PERSON>0 destroy-44 surface.n.01 Ø Ø

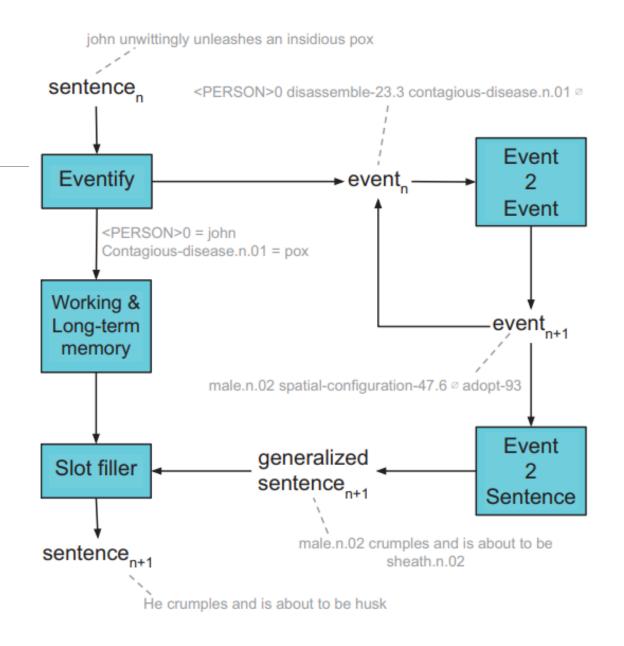
Review: Story Realization

Extract events from stories

Generate the plot using a seq2seq network

Use an ensemble of methods to find the best sentence given an event

Get a confidence score from each model, and accept the sentence if it's above a threshold



Review: Story Cloze Test

Gina was worried the cookie dough in the tube would be gross.

She was very happy to find she was wrong.

The cookies from the tube were as good as from scratch.

Gina intended to only eat 2 cookies and save the rest.

A. Gina liked the cookies so much she ate them all in one sitting.



B. Gina gave the cookies away at her church.

14

CMSC 491/691: Interactive Fiction and Text Generation

Search and Planning

AIMA Chapters 3 and 7



Learning Objectives

Remember how to setup a search problem

Review basic types of tree search algorithms

Define & implement a search problem (for Action Castle)

Problem-Solving Agents

A problem-solving agent must **plan**.

The computational process that it undertakes is called **search**.

It will consider a **sequence of actions** that form a **path** to a **goal state**.

Such a sequence is called a **solution**.

3	go south
J •	go boach
4.	catch fish with po
5.	go north
6.	pick rose
7.	go north
8.	go up
9.	get branch
10.	go down
11.	go east
12.	give the troll
	the fish

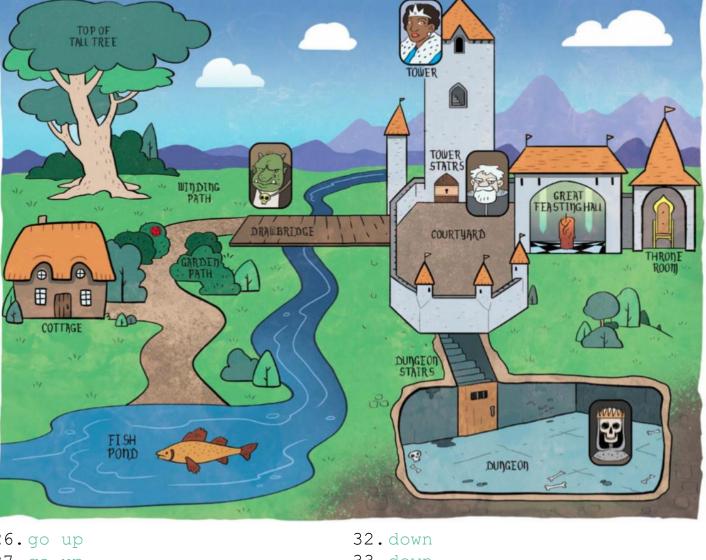
1. take pole

2. go out

13.go east
14.hit guard with branch
15.get key
ble 16.go east
17.get candle
18.go west
19.go down
20.light lamp
21.go down
22.light candle

23.read runes
24.get crown

25.go up



26.go up	32.down
27.go up	33.down
28.unlock door	34.east
29.go up	35.east
30. give rose to the princess	36.wear crown
31. propose to the princess	37. sit on throne

Review of Search Problems

AIMA 3.1-3.3

Formal Definition of a Search Problem

- 1. States: a set S
- 2. An **initial state** s_i∈ S
- 3. Actions: a set A

∀ s Actions(s) = the set of actions that can be executed in **s**.

4. Transition Model: ∀ s∀ a∈Actions(s)

Result(s, a)
$$\rightarrow$$
 s_r

 \mathbf{s}_{r} is called a successor of \mathbf{s}

- 5. Path cost (Performance Measure):
 Must be additive, e.g. sum of distances,
 number of actions executed, ...
 - c(x,a,y) is the step cost, assumed ≥ 0
 - (where action a goes from state x to state y)
- Goal test: Goal(s)

s is a goal state if **Goal(s)** is true. Can be implicit, e.g. **checkmate(s)**

States: A state of the world says which objects are in which cells.

In a simple two cell version,

- the agent can be in one cell at a time
- each cell can have dirt or not

2 positions for agent * 2^2 possibilities for dirt = 8 states.

With n cells, there are $n*2^n$ states.



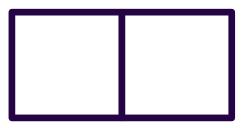
States: A state of the world says which objects are in which cells.

In a simple two cell version,

- the agent can be in one cell at a time
- each cell can have dirt or not

2 positions for agent * 2^2 possibilities for dirt = 8 states.

With n cells, there are $n*2^n$ states.









States: A state of the world says which objects are in which cells.

In a simple two cell version,

- the agent can be in one cell at a time
- each cell can have dirt or not

2 positions for agent * 2^2 possibilities for dirt = 8 states.

With n cells, there are $n*2^n$ states.

One state is designated as the **initial state**

Goal states: States where everything is clean.





















37



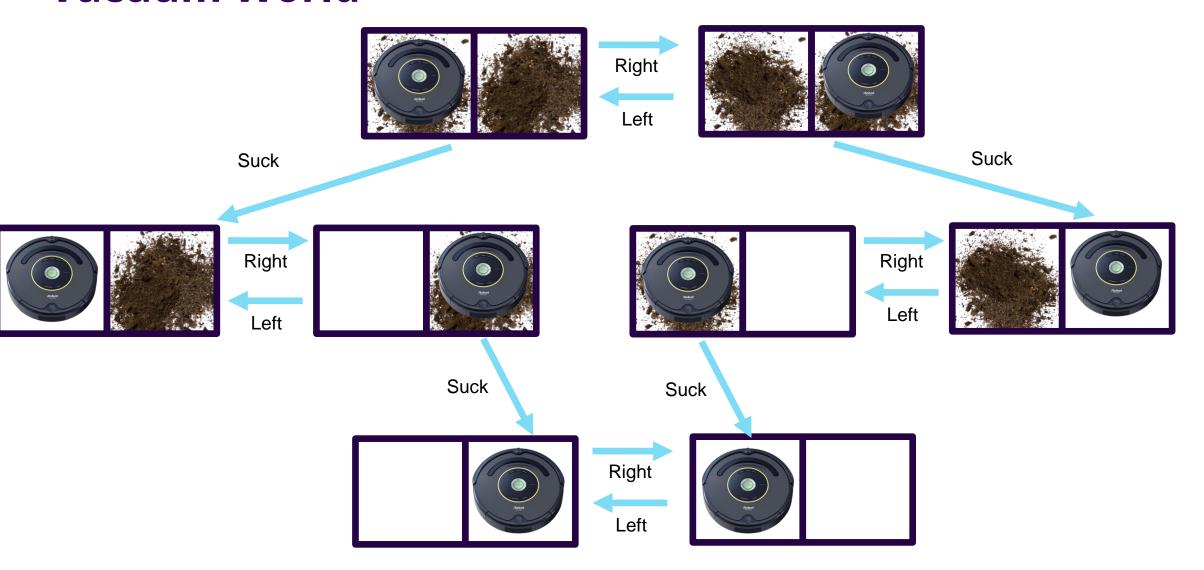
Actions:

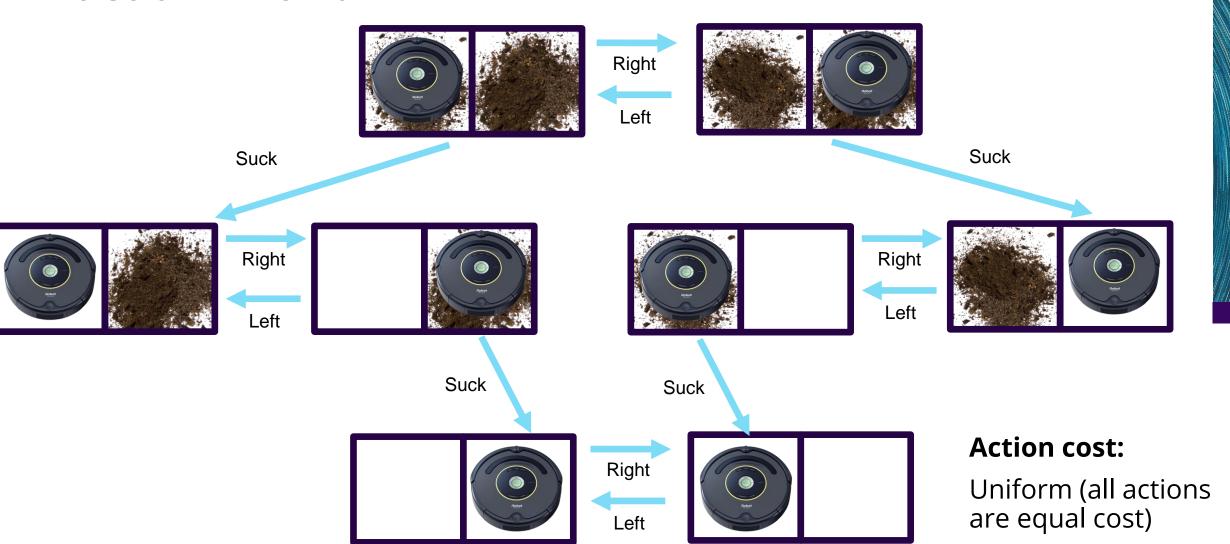
- Suck
- Move Left
- Move Right
- (Move Up)
- (Move Down)

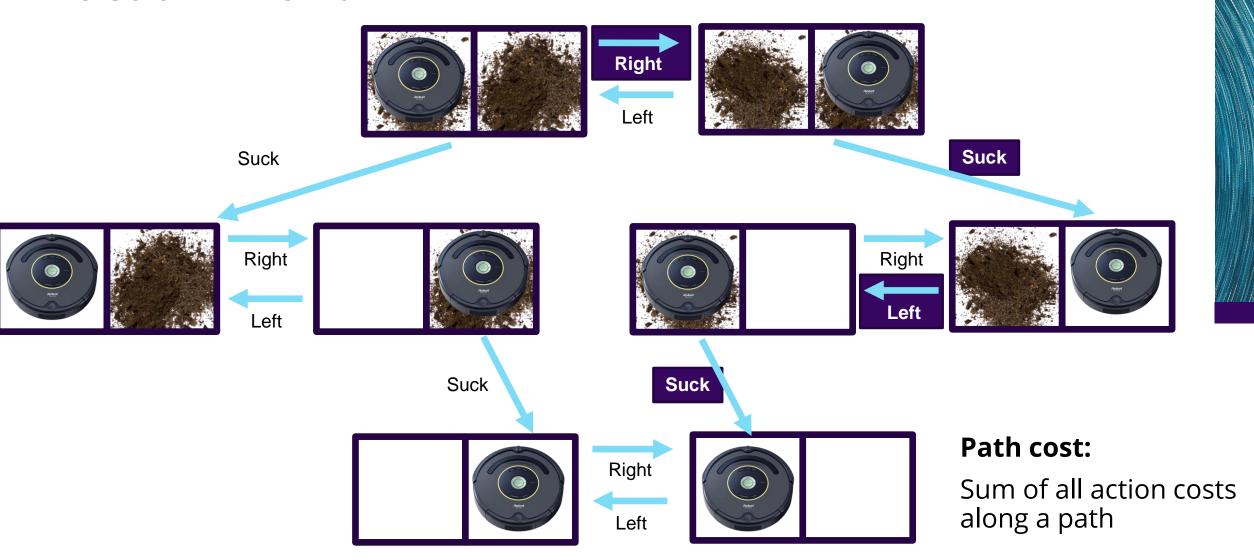
Transition:

Suck – removes dirt

Move – moves in that direction, unless agent hits a wall, in which case it stays put.







Initial state







Suck

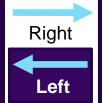
Suck



Right Left





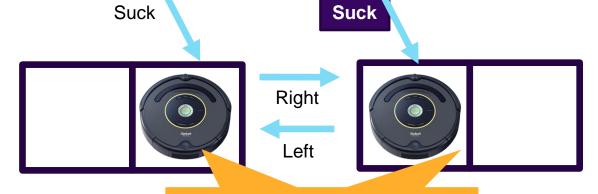






Solution:

A path from the initial state to a goal state



Goal states

Search Algorithms

Useful Concepts

State space: the set of all states reachable from the initial state by *any* sequence of actions

- When several operators can apply to each state, this gets large very quickly
- Might be a proper subset of the set of configurations

Path: a sequence of actions leading from one state s_j to another state s_k

Solution: a path from the initial state s_i to a state s_f that satisfies the goal test

Search tree: a way of representing the paths that a search algorithm has explored. The root is the initial state, leaves of the tree are successor states.

Frontier: those states that are available for **expanding** (for applying legal actions to)

Solutions and *Optimal* Solutions

A *solution* is a sequence of actions from the initial state to a goal state.

Optimal Solution: A solution is optimal if no solution has a lower path cost.

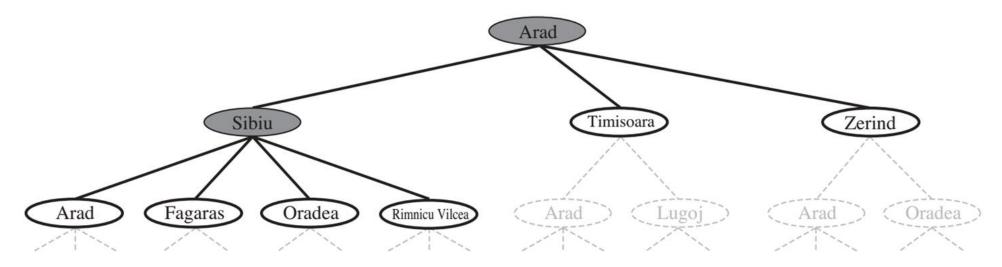
Basic search algorithms: Tree Search

Generalized algorithm to solve search problems

Enumerate in some order all possible paths from the initial state

- Here: search through explicit tree generation
 - ROOT= initial state.
 - Nodes in search tree generated through transition model
 - Tree search treats different paths to the same node as distinct

Generalized tree search



function TREE-SEARCH(*problem, strategy*) return a solution or failure

Initialize frontier to the *initial state* of the *problem* do

The strategy determines search process!

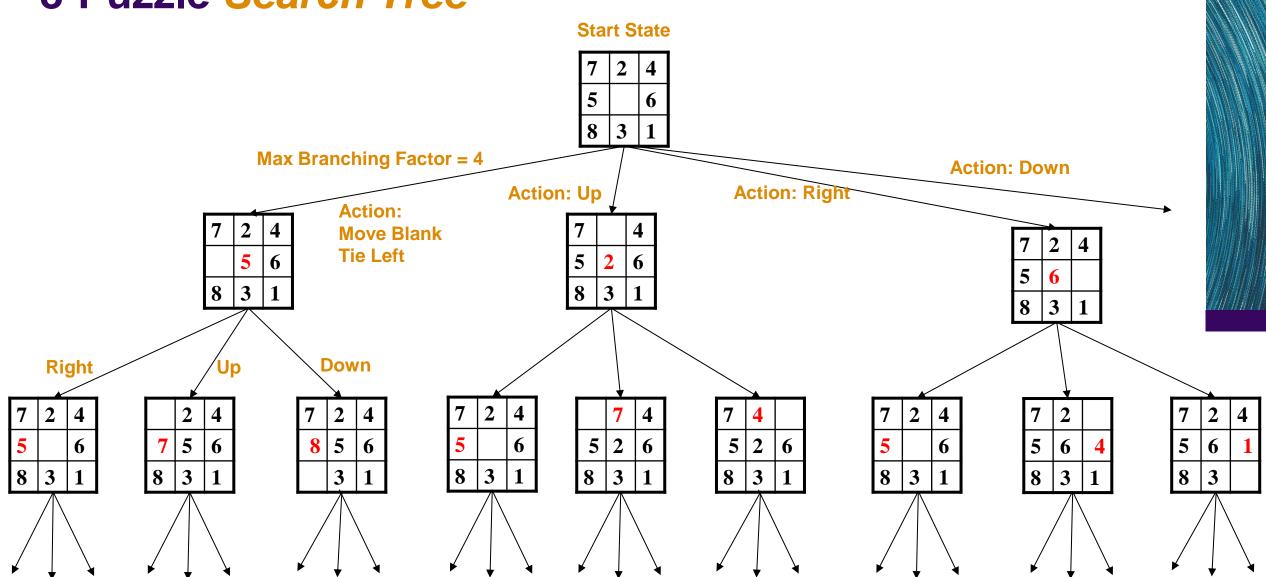
if the frontier is empty then return failure

choose leaf node for expansion according to strategy & remove from frontier

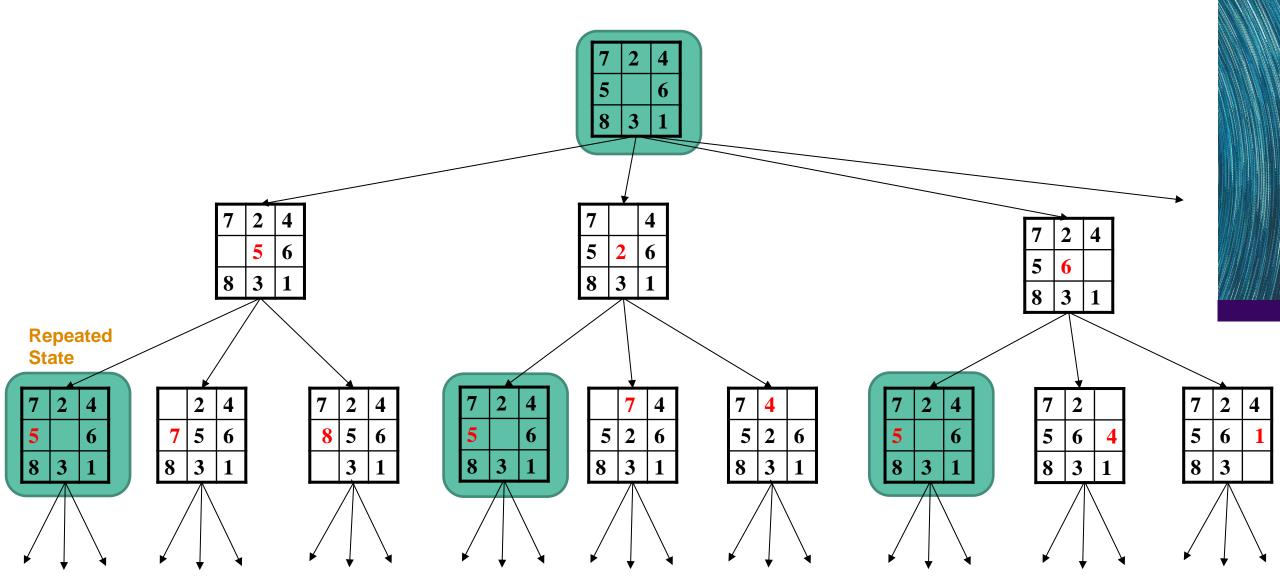
if node contains goal state then return solution

else expand the node and add resulting nodes to the frontier

8-Puzzle Search Tree



8-Puzzle Search Tree



Graph Search vs Tree Search

function Tree-Search(*problem*) **returns** a solution, or failure initialize the frontier using the initial state of *problem* **loop do**

if the frontier is empty then return failurechoose a leaf nose and remove it from the frontierif the node contains a goal state then return the corresponding solution expand the chosen node, adding the resulting nodes to the frontier

function Graph-Search(*problem*) returns a solution, or failure initialize the frontier using the initial state of *problem initialize the explored set to be empty*

loop do

if the frontier is empty **then return** failure choose a leaf node and remove it from the frontier

if the node contains a goal state then return the corresponding solution

add node to the explored set

expand the chosen node, adding the resulting nodes to the frontier only if not in the frontier of explored set

Search Strategies

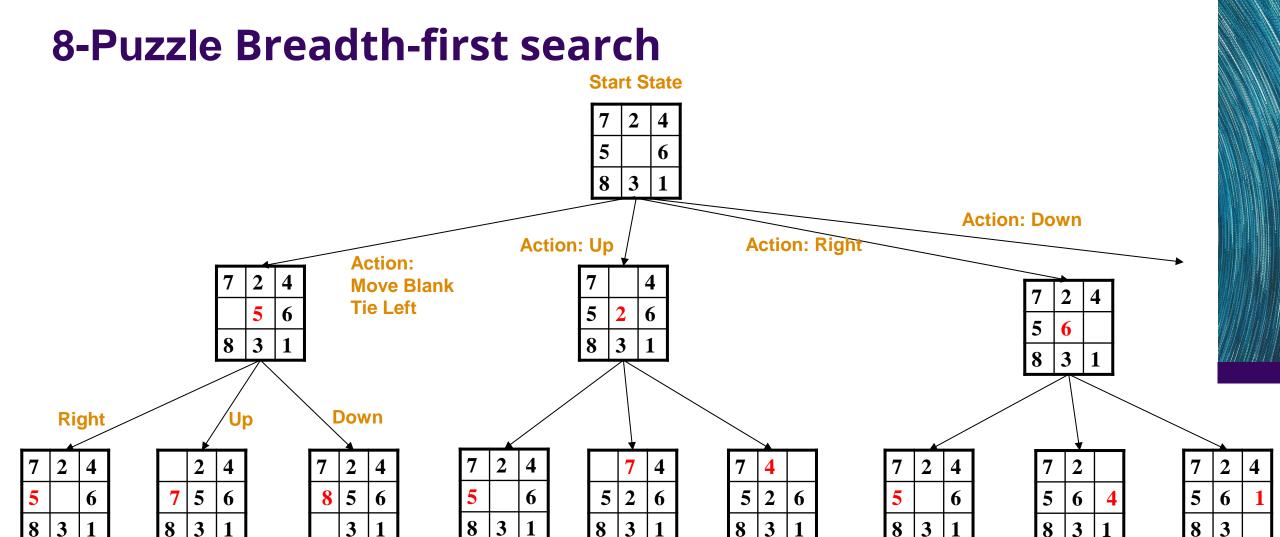
Several classic search algorithms differ only by the order of how they expand their search trees

You can implement them by using different queue data structures

Depth-first search = LIFO queue

Breadth-first search = FIFO queue

Greedy best-first search or A* search = Priority queue



Module 2 Review

Search Algorithms

Dimensions for evaluation

- Completeness- always find the solution?
- Optimality finds a least cost solution (lowest path cost) first?
- Time complexity # of nodes generated (worst case)
- Space complexity # of nodes simultaneously in memory (worst case)

Time/space complexity variables

- *b, maximum branching factor* of search tree
- *d, depth* of the shallowest goal node
- *m*, *maximum length* of any path in the state space (potentially ∞)

Properties of breadth-first search

Complete? Yes (if *b* is finite)

Optimal? Yes, if cost = 1 per step

(not optimal in general)

Time Complexity? $1+b+b^2+b^3+...+b^d = O(b^d)$

Space Complexity? $O(b^d)$ (keeps every node in memory)

Time/space complexity variables

- *b, maximum branching factor* of search tree
- *d, depth* of the shallowest goal node
- *m*, *maximum length* of any path in the state space (potentially ∞)

BFS versus DFS

Breadth-first

- ☑ Complete,
- ✓ Optimal
- **☑** but uses $O(b^d)$ space

Time/space complexity variables

- b, maximum branching factor of search tree
- d, depth of the shallowest goal node
- *m*, *maximum length* of any path in the state space (potentially ∞)

Depth-first

- Not complete unless m is bounded
- Not optimal
- \blacksquare Uses $O(b^m)$ time; terrible if m >> d
- **☑** but only uses O(**b*m) space**

Exponential Space (and time) Is Not Good...

- Exponential complexity uninformed search problems *cannot* be solved for any but the smallest instances.
- (Memory requirements are a bigger problem than execution time.)

DEPTH	NODES	TIME	MEMORY
2	110	0.11 milliseconds	107 kilobytes
4	11110	11 milliseconds	10.6 megabytes
6	10^6	1.1 seconds	1 gigabytes
8	10^8	2 minutes	103 gigabytes
10	10^{10}	3 hours	10 terabytes
12	10^{12}	13 days	1 petabytes
14	10^{14}	3.5 years	99 petabytles

Assumes b=10, 1M nodes/sec, 1000 bytes/node

Action Castle

Art: Formulating a Search Problem

Decide:

Which properties matter & how to represent

Initial State, Goal State, Possible Intermediate States

Which actions are possible & how to represent

Operator Set: Actions and Transition Model

Which action is next

Path Cost Function

Formulation greatly affects combinatorics of search space and therefore speed of search

Action Castle Map Navigation

Let's consider the sub-task of navigating from one location to another.

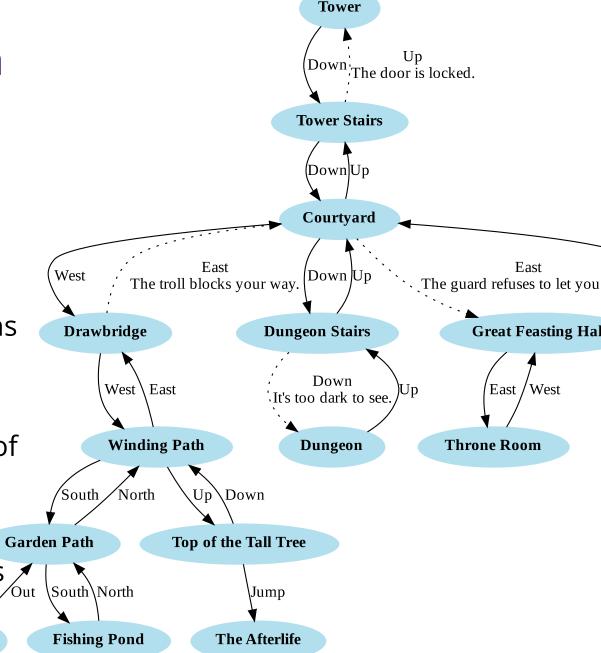
Formulate the *search problem*

- States: locations in the game
- Actions: move between connected locations
- Goal: move to a particular location like the Throne Room
- Performance measure: minimize number of moves to arrive at the goal

Find a **solution**

Algorithm that returns sequence of actions to get from the start sate to the goal.

Cottage



Module 2 Review

59

```
def BFS(game, goal_conditions):
 command_sequence = []
 if goal_test(game, goal_conditions): return 
 frontier = queue.Queue()
 frontier.put((game, command_sequence))
                                                       prevents us from
 visited = dict()
                                                        revising states.
 visited[get_state(game)] = True
 while not frontier.empty():
    (current_game, command_sequence) = frontier.get()
   current_state = get_state(current_game)
   parser = Parser(current_game)
   available_actions = get_available_actions(current_game)
   for command in available_actions:
     # Clone the current game with its state
     new_game = copy.deepcopy(current_game)
     # Apply the command to it to get the resulting state
      parser = Parser(new_game)
     parser.parse_command(command)
     new_state = get_state(new_game)
     # Update the sequence of actions that we took to get to
     new_command_sequence = copy.copy(command_sequence)
     new_command_sequence.append(command)
      if not new_state in visited:
       visited[new_state] = True
       if goal test(new_game, goal_conditions):
        frontier.put((new_game, new_command_sequence))
  # Return None to indicate there is no solution.
  return None
```

13

15

16 17

18 19

20

21

23

24

25

26

27

28

30

32

The frontier tracks order of unexpanded search nodes. Here we're using a FIFO queue

The visited dictionary

TODO: implement get_state()

get_available_actions() to return all commands that could be used here.

TODO: implement get_available_actions()

The parser can execute this command to get the resulting state.

Check to see if this state satisfies the goal test, if so, return the command sequence that got us here.

TODO: implement goal_test()

```
def BFS(game, goal_conditions):
     command_sequence = []
     if goal_test(game, goal_conditions): return command_sequence
                                                           Tip: We can store multiple objects
     frontier = queue.Queue()
     frontier.put((game, command_sequence))
                                                                on the frontier as a tuple.
     visited = dict()
     visited[get_state(game)] = True
                                                      Tip: To be used a key in the dictionary get_state()
10
11
     while not frontier.empty():
                                                                must return an immutable object
       (current_game, command_sequence) = frontier.get
13
       current_state = get_state(current_game)
       parser = Parser(current_game)
14
15
       available_actions = get_available_actions(current_game)
16
17
       for command in available_actions:
         # Clone the current game with its state
                                                        Tip: use deepcopy here
18
19
         new game = copy.deepcopy(current_game)
20
         # Apply the command to it to get the resulting state
21
         parser = Parser(new_game)
         parser.parse_command(command)
23
         new_state = get_state(new_game)
24
         # Update the sequence of actions that we took to get to the resulting state
25
         new_command_sequence = copy.copy(command_sequence)
26
         new_command_sequence.append(command)
                                                             Tip: For BFS, apply the goal test before
27
         if not new_state in visited:
28
           visited[new_state] = True
                                                              putting the new item on the frontier
           if goal test(new game, goal conditions):
           frontier.put((new_game, new_command_sequence))
30
31
     # Return None to indicate there is no solution.
32
     return None
```

Action Castle

Let's consider the full game.

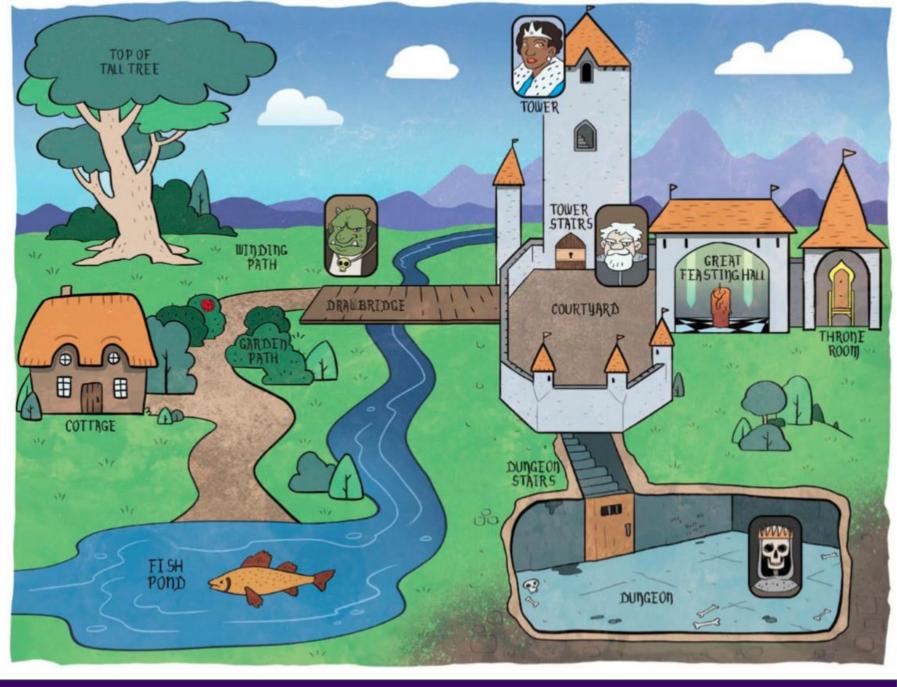
Actions

Start State

Transitions

State Space

Goal test



Actions

Go

Move to a location

Get

Add an item to inventory

Special

Perform a special action with an item like "Catch fish with pole"

Drop

Leave an item in current location



State Info

Location of Player
Items in their inventory
Location of all items /
NPCs

Blocks like

- Troll guarding bridge,
- Locked door to tower,
- Guard barring entry to castle



In-Class Activity

https://laramartin.net/interactive-fiction-class//in_class_activities/search/action-castle-search.html

https://bit.ly/4eSjws8

Classical Planning

AIMA Chapter 11

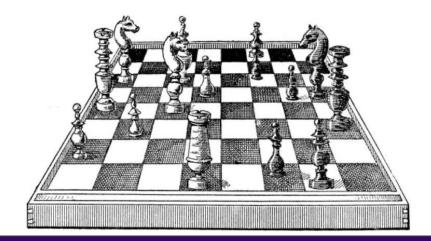
Classical Planning

The task of finding a sequence of action to accomplish a goal in a deterministic, fully-observable, discrete, static environment.

If an environment is:

- Deterministic
- Fully observable

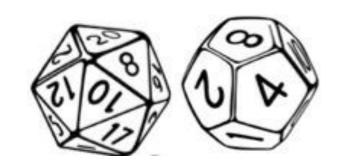
The solution to any problem in such an environment is a fixed sequence of actions.



In environments that are

- Nondeterministic or
- Partially observable

The solution must recommend different future actions depending on the what percepts it receives. This could be in the form of a *branching strategy*.



Representation Language

Planning Domain Definition Language (PDDL) express actions as a schema

```
Action name
                                                                  Variables
              (:action go
               :parameters (?dir - direction
                                 ?p - player
                                 ?l1 - location ?l2 - location)
               :precondition (and
                                                                      Preconditions
                                       (at?p?|1)
  Preconditions and effects are
                                       (connected ?l1 ?dir ?l2)
conjunctions of logical sentences
                                       (not (blocked ?l1 ?dir ?l2)))
               :effect (and
                                       (at?p?l2)
                                                                           Effects
                                       (not (at ?p ?[1)))
                            These logical sentences are literals –
                            positive or negated atomic sentences
```

State Representation

In PDDL, a **state** is represented as a **conjunction** of logical sentences that are

ground atomic fluents. PDDL uses database semantics.

Ground means they contain no variables Atomic sentences contain just a single predicate

of the world that can change over time.

Closed world assumption. Any fluent not mentioned is false. Unique names are distinct.

Action Schema has variables

```
(:action go
:parameters (?dir - direction ?p - player ?l1 - location ?l2 - location)
:precondition (and (at ?p ?l1) (connected ?l1 ?dir ?l2) (not (blocked ?l1 ?dir ?l2)))
:effect (and (at ?p ?l2) (not (at ?p ?l1)))
)
```

State Representation arguments are constants fluents may change over time

```
(connected cottage out gardenpath)
(connected gardenpath in cottage)
(connected gardenpath south fishingpond)
(connected fishingpond north gardenpath)
(at npc cottage)
```

Successor States

A **ground action** is **applicable** if if every positive literal in the precondition is true, and every negative literal in the precondition is false

Ground Action no variables

```
(:action go
:parameters (out, npc, cottage, gardenpath)
:precondition (and (at npc cottage) (connected cottage out gardenpath)

(not (blocked cottage out gardenpath)))
:effect (and (at npc gardenpath) (not (at npc cottage)))
)
```

Initial State

(connected cottage out gardenpath)
(connected gardenpath in cottage)
(connected gardenpath south fishingpond)
(connected fishingpond north gardenpath)
(at npc cottage)

Negative literals in the effects are kept in a **delete list** DEL(), and positive literals are kept in an **add list** ADD()

Result

New state reflecting the effect of applying the ground action

```
(connected cottage out gardenpath)
(connected gardenpath in cottage)
(connected gardenpath south fishingpond)
(connected fishingpond north gardenpath)
(at npc gardenpath)
```

Domain

Set of Action Schema

Module 2 Review

```
(define (domain action-castle)
 (:requirements :strips :typing)
 (:types player location direction it
 (:action go
   :parameters (?dir - direction ?p - player
          ?l1 - location ?l2 - location)
   :precondition (and (at?p?l1)
              (connected ?l1 ?dir ?l2)
              (not (blocked ?l1 ?dir ?l2)))
   :effect (and (at ?p ?l2) (not (at ?p ?l1)))
 (:action get
   :parameters (?item - item
          ?p - player
          ?I1 - location)
   :precondition (and (at?p?l1)
              (at ?item ?l1))
   :effect (and (inventory ?p ?item)
          (not (at ?item ?l1)))
```

Problem

```
(define (problem navigate-to-location)
 (:domain action-castle)
 (:objects
   npc - player
   cottage gardenpath fishingpond gardenpath
    windingpath talltree drawbridge courtyard
    towerstairs tower dungeonstairs dungeon
    greatfeastinghall throneroom - location
   in out north south east west up down - direction
                                                   Initial State
 (:init
   (at npc cottage)
   (connected cottage out gardenpath)
   (connected gardenpath in cottage)
   (connected gardenpath south fishingpond)
   (connected fishingpond north gardenpath)
 (:goal (and (at npc throneroom)))
```

Goal

73

Algorithms for Classical Planning

We can apply **BFS** to the **initial state** through possible states looking for a **goal**.

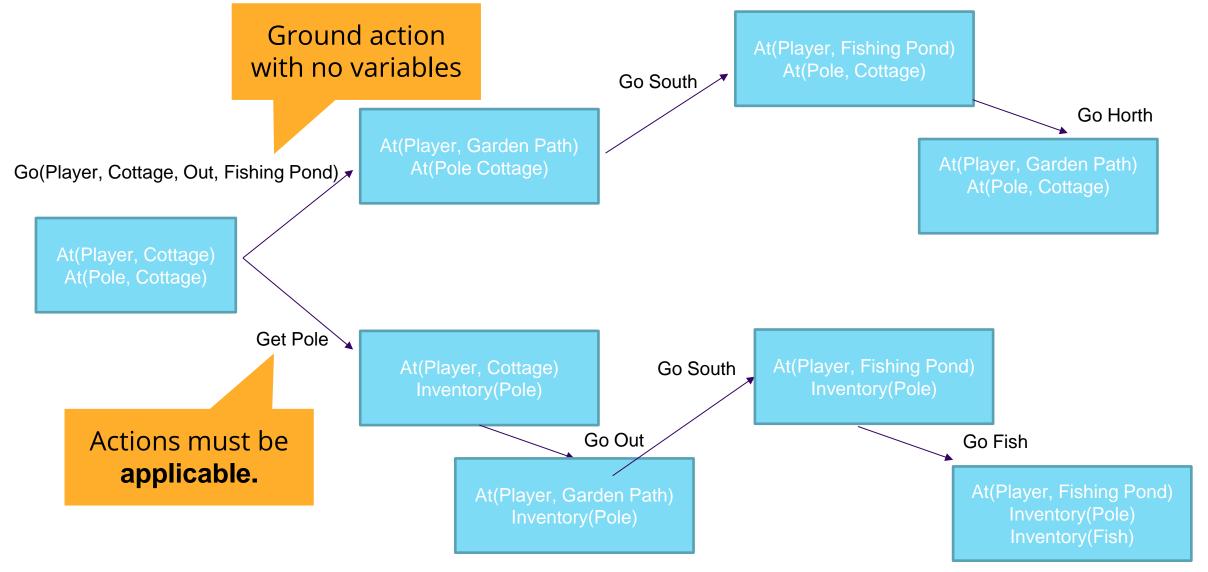
An advantage of the **declarative representation** of action schemas is that we can also **search backwards**.

Start with a goal and work backwards towards the initial state.



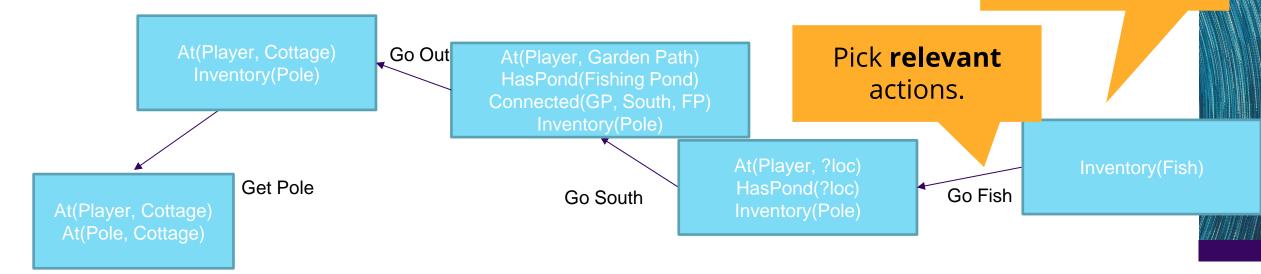
In our Action Castle example, this would help us with the branching problem that the **drop** action introduced. If we work backwards from the goal, then we realize that we don't ever need to drop an item for the correct solution.

Forward State-Space Search for Planning



Backward State-Space Search for Planning aka Regression Search

Start with the goal, work backwards to initial state



Given a goal **g** and action **a**, the **regression** from g to a gives a state **g**' description whose literals are given by:

POS(g') = (POS(g)-ADD(a)) U POS(Preconditions(a))

NEG(g') = (NEG(g)-DF(a)) U NEG(Preconditions(a))

Negative literals in the effects are kept in a delete list DEL

Positive literals in the effects are kepter www.ADD list

Heuristics for Planning

Neither forward nor backward search is efficient without good heuristics.

In search a heuristic function h(s) estimates the distance from a state to the goal.

Admissible heuristics never over-estimate the true distance, and can be used with **A* search** to find optimal solutions.

Admissible heuristics can be derived from a **relaxed problem** that is easier to solve. For the **ignore preconditions heuristic** relaxes the problem.

Hierarchical Planning

Instead of using atomic actions, we can define actions at **higher levels of abstraction**.

Hierarchical decomposition organizes actions into high-level functions, composed of more fine-grained function, composed of atomic actions.

Plan out sequence of high level actions, reclusively **refine the plan** until we've got atomic actions.

Tricky to ensure that the resulting plan is optimal.

Planning and Games

Planning can be used for AI characters

In our current text adventure games, all of the non-player characters are boring!

- Why doesn't the princess try to escape the tower and claim the throne herself?
- Why doesn't the troll come hunting for food and eat us or the guard?
- Why is the ghost of the king stuck in the dungeon?

We could give each of them goals, and have them try to plan out and play the game alongside the player.

Generating Puzzles

In HW2, we were able to generate descriptions of locations and items and automatically incorporate them into our text adventure games.

Could we use planning to automatically generate:

- 1. Puzzles?
- 2. Special actions?

Let's say a player needs a **sword** and we decide to make the game more challenging by not putting one anywhere in the game.

Could we generate an action that results in the creation of a sword?

Action: forge a sword

Effects: a sword is created

Preconditions: molten metal, a cast of a sword, an anvil, a hammer